

Interest Points

Computer Vision (CS0029)

Outline

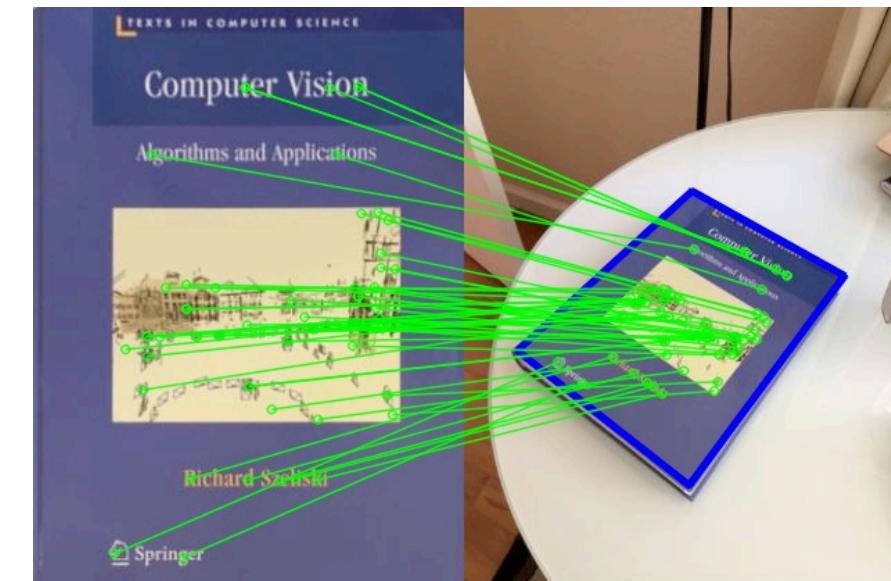
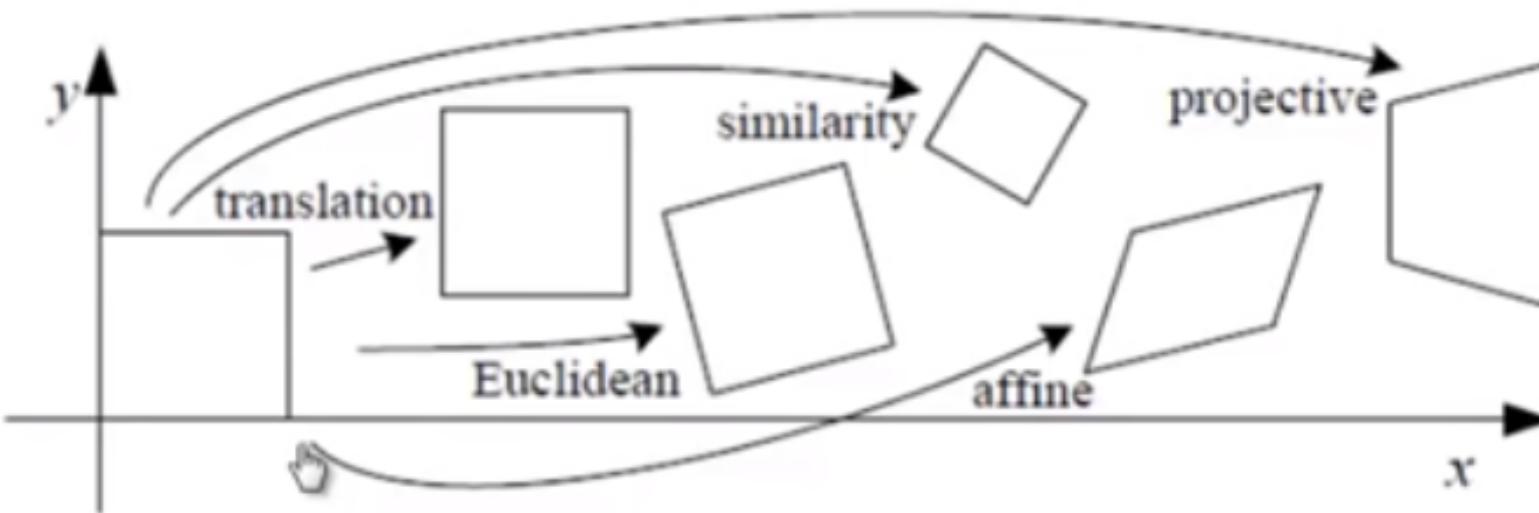
- **Local Invariant Features**
 - Motivation
 - Requirements, Invariances
- Keypoint Localization
 - Feature from Accelerated Segment Test (FAST)
 - Harris
 - Shi-Tomasi
- Scale Invariant Region Selection
 - Automatic scale selection
 - Laplacian-of-Gaussian detector
 - Difference-of-Gaussians detector
- Local Descriptors
 - Orientation normalization
 - SIFT

Consider Image Matching



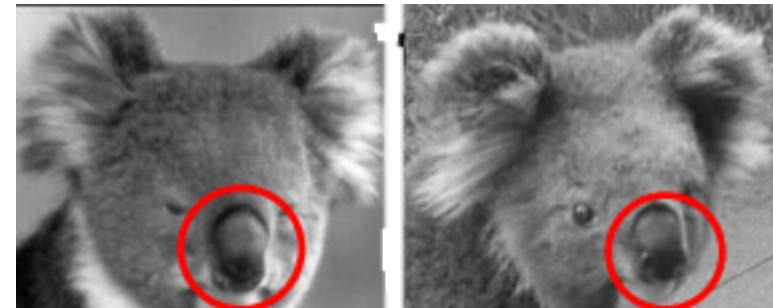
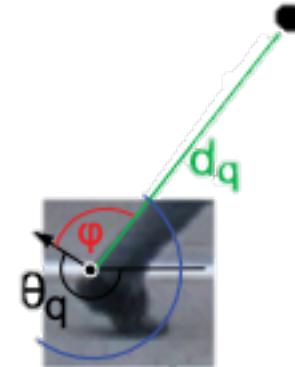
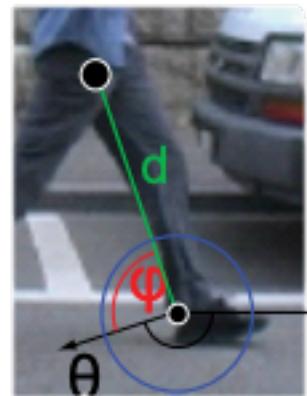
Consider Image Matching

- Estimate the transformation between two images
 - We have to find the matching point between two images first
 - We need to find local features



Global vs. Local Matching

- Global image representation can be easily corrupted (e.g., by even small occlusion)
- Instead, describe and match using multiple local regions
- Increased robustness to occlusions, articulation, within-category variation



Goal of Local Features

- Find points in an image that can be:
 - Found in other images
 - Found precisely (repeatability) – well localized
 - Found reliably – well matched
- Want to compute a fundamental matrix to recover geometry
- Robotics/vision: see how a bunch of points move from one frame to another. Allows computation of how camera moved -> depth -> moving objects
 - E.g. structure from motion, <https://www.youtube.com/watch?v=i7ierVkXYa8>
- Panorama

Application: Panorama



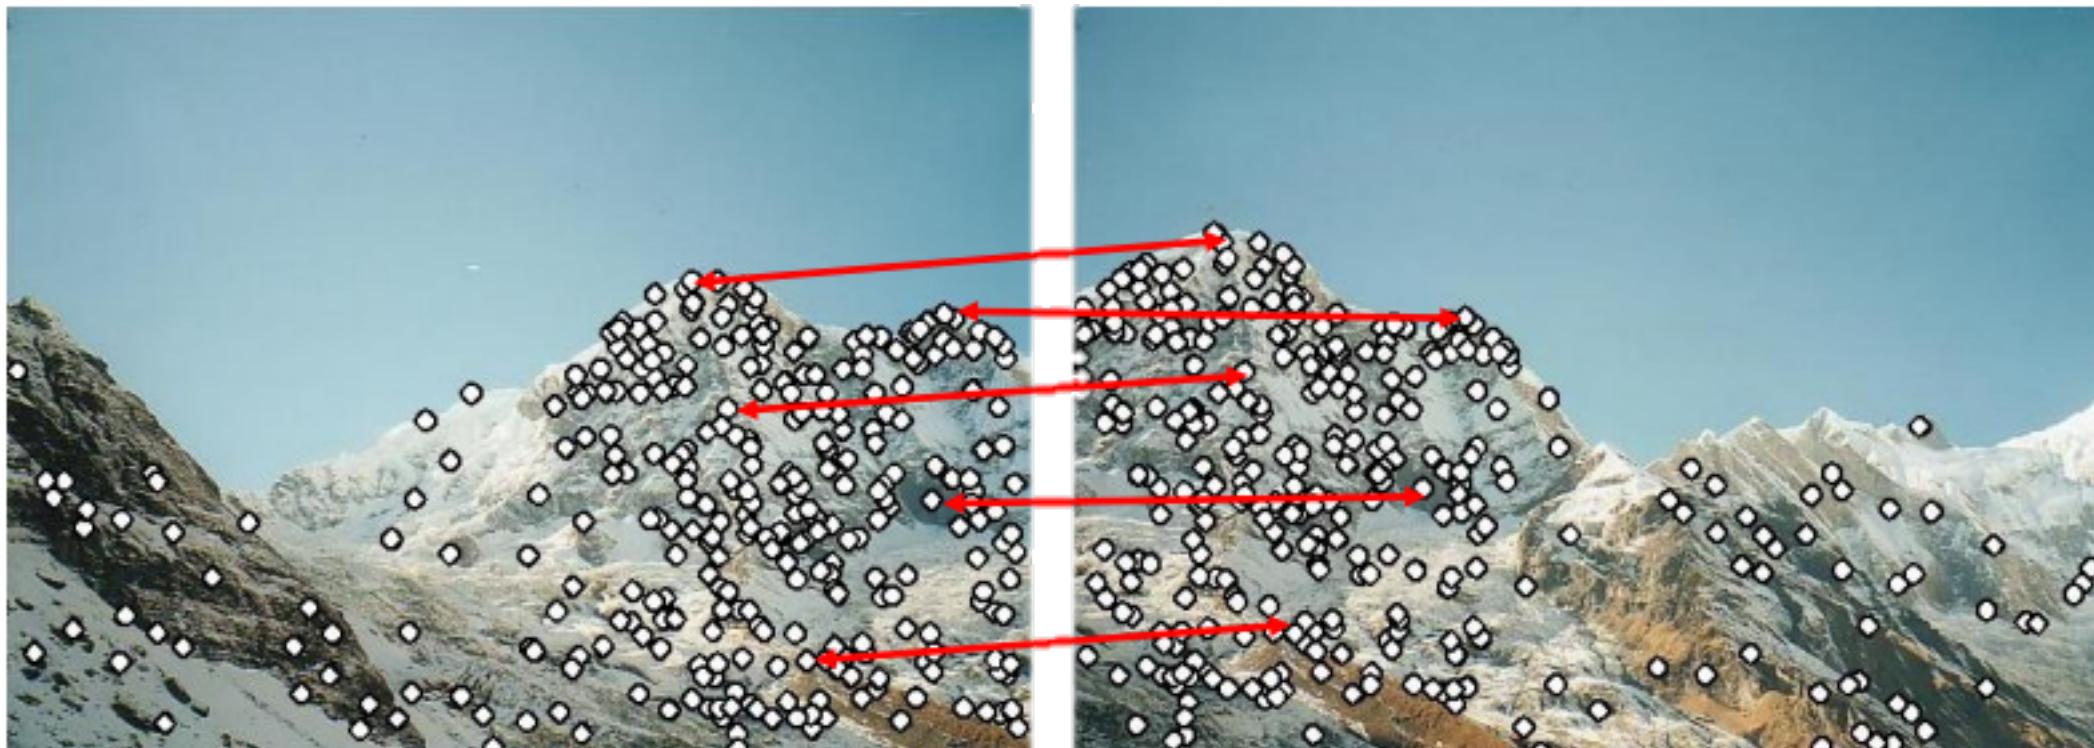
Application: Panorama

Detect features (feature points) in both images



Application: Panorama

Match: find corresponding pairs

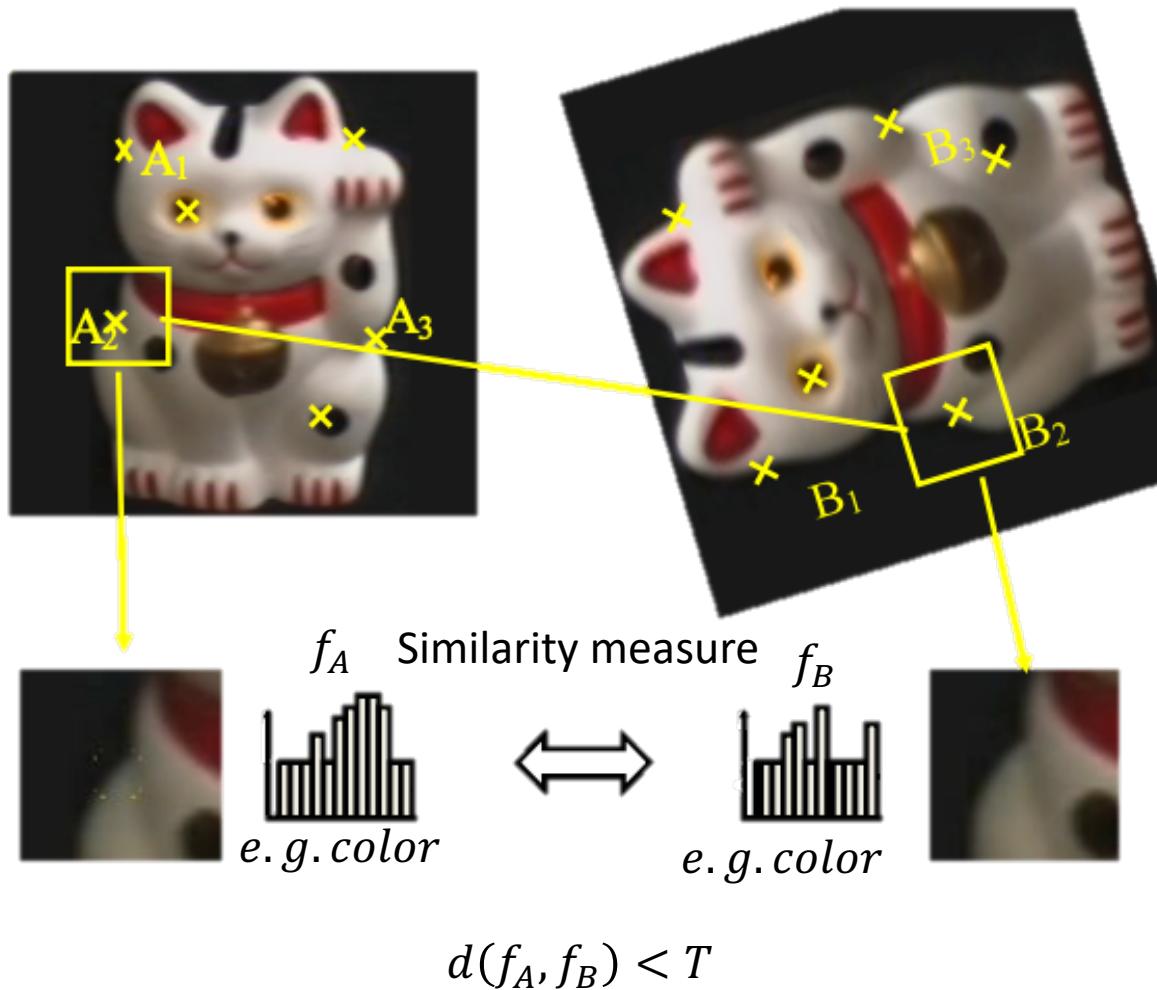


Application: Panorama

Use the pairs to transform/align images



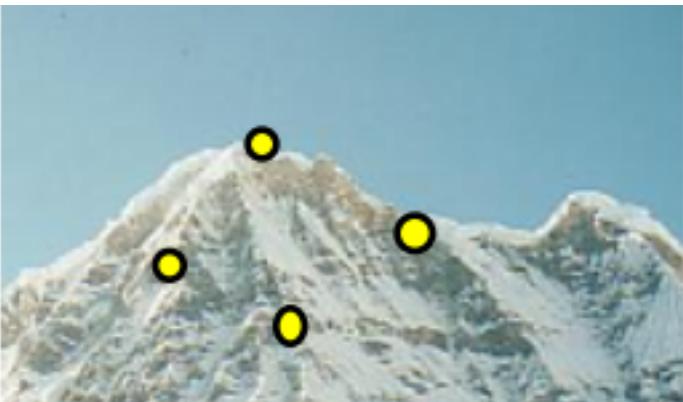
General Approach



- 1. Find a set of distinctive keypoints
- Define a region around each keypoint
- Extract the region content
- Compute a local descriptor from the region
- Match local descriptors

Common Requirements

- Problem 1:
 - Detect the same points independently in both images

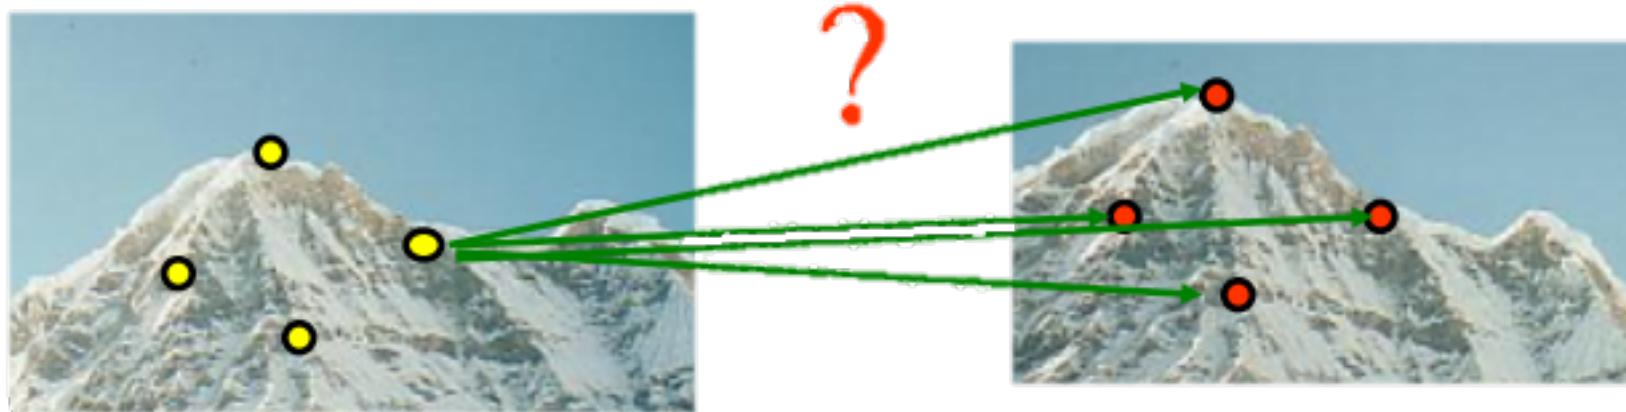


No chance to match!

We need a repeatable detector!

Common Requirements

- Problem 1:
 - Detect the same points independently in both images
- Problem 2:
 - For each point, correctly identify the corresponding point



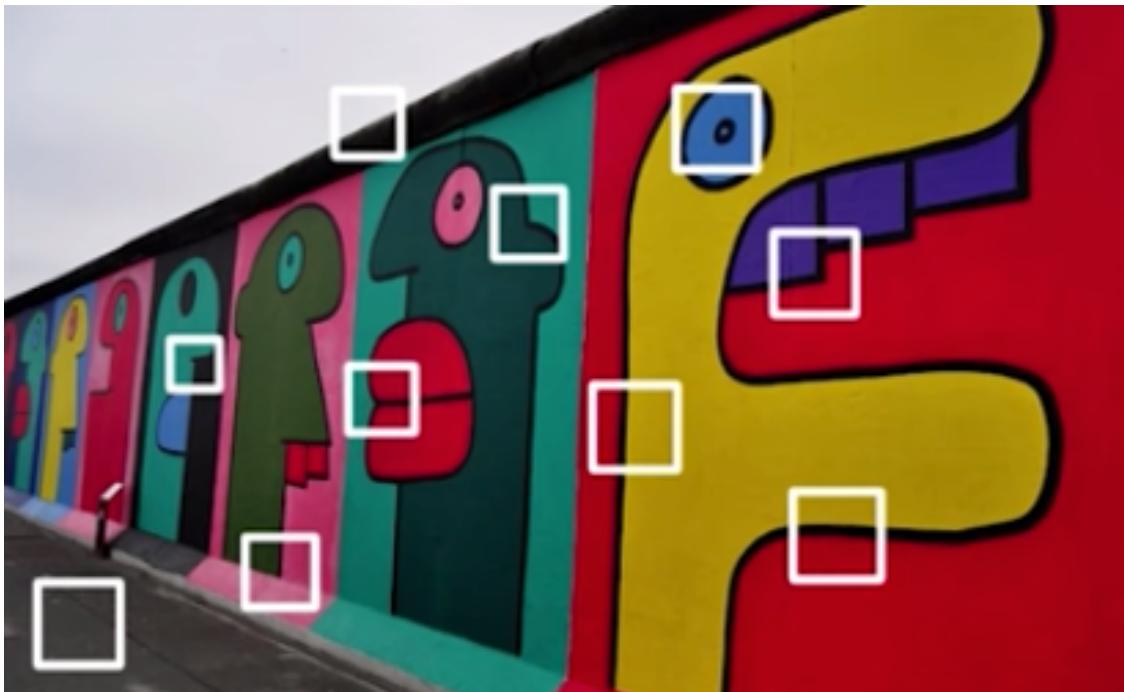
We need a reliable and distinctive descriptor

More Motivation

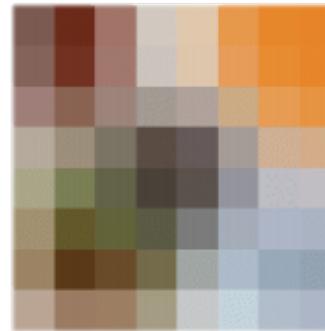
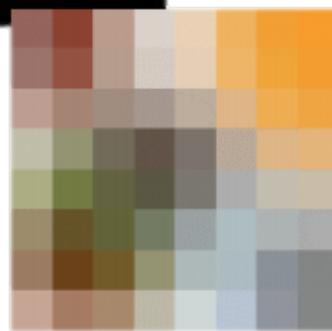
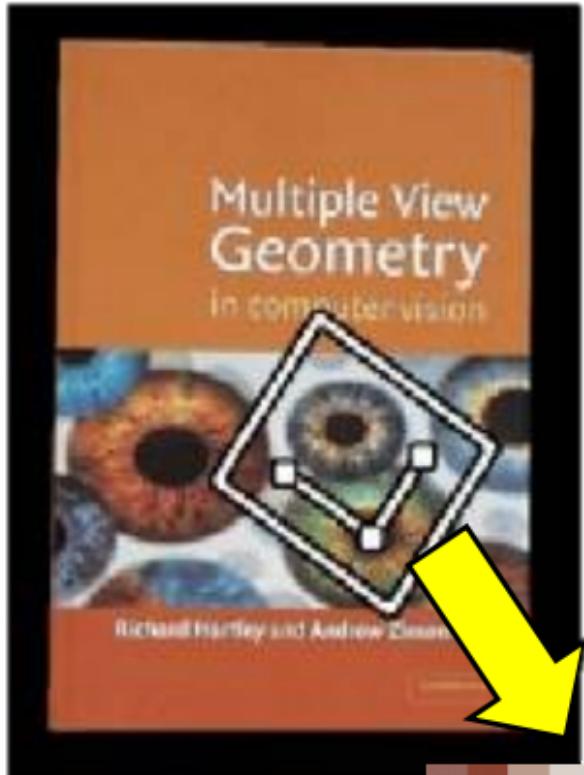
- Feature points are used also for:
 - Image alignment
 - 3D reconstruction
 - Motion tracking
 - Object recognition
 - Indexing and database retrieval
 - Robot navigation
 -

Practice

- Select the ones that can serve as good features



Invariance: Geometric Transformations

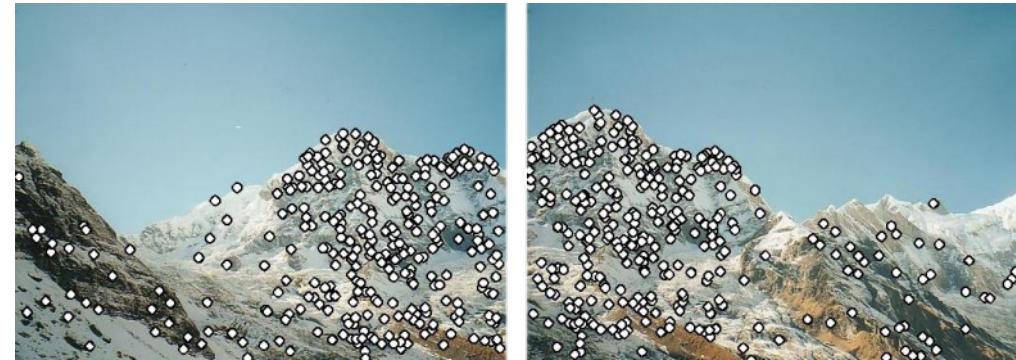


Invariance: Photometric Transformations



Key Takeaways

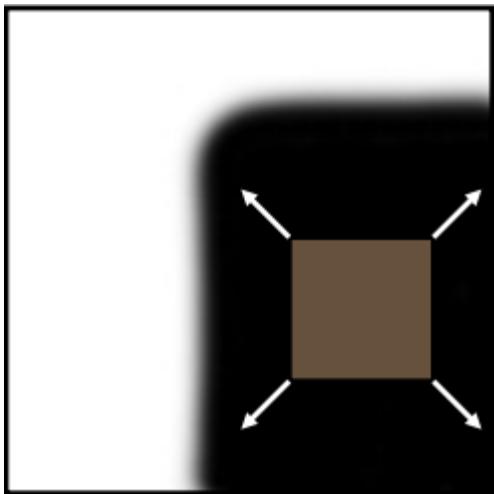
- Repeatable/precision:
 - The same feature can be found in several images despite geometric and photometric transformation
- Saliency/matchability
 - Each feature has a distinctive description
- Compactness and efficiency
 - Many fewer features than image pixels
 - close to real-time performance
- Locality
 - A feature occupies a relatively small area of the image, robust to clutter and occlusion
- Invariance
 - Invariant to translation, rotation, scale changes, and other simple (affine) transformations
 - Robust to lighting variations, noise, blur, quantization



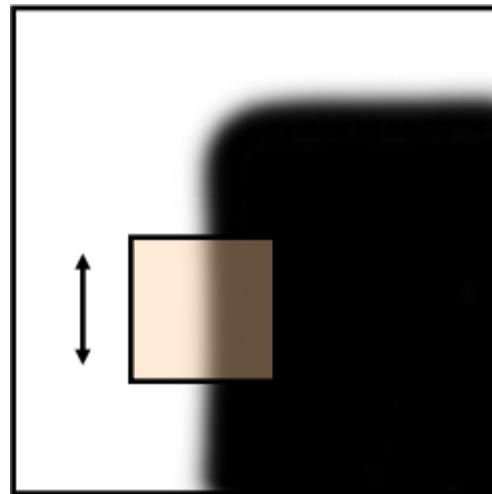
Outline

- Local Invariant Features
 - Motivation
 - Requirements, Invariances
- Keypoint Localization
 - Feature from Accelerated Segment Test (FAST)
 - Harris
 - Shi-Tomasi
- Scale Invariant Region Selection
 - Automatic scale selection
 - Laplacian-of-Gaussian detector
 - Difference-of-Gaussians detector
- Local Descriptors
 - Orientation normalization
 - SIFT

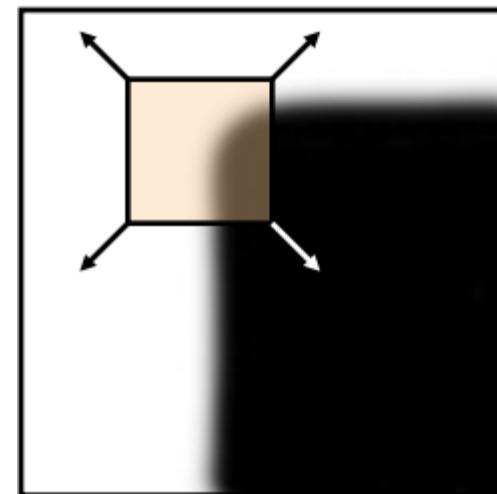
Flat vs. Edge vs. Corner (Precision)



“flat” region: no change
in all directions



“edge”: no change
along the edge direction



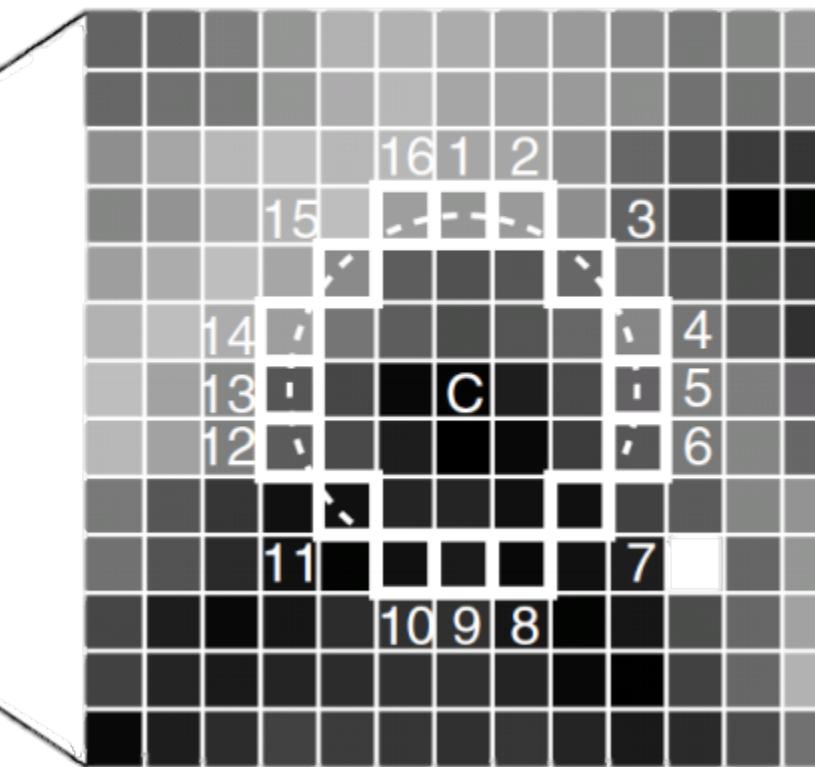
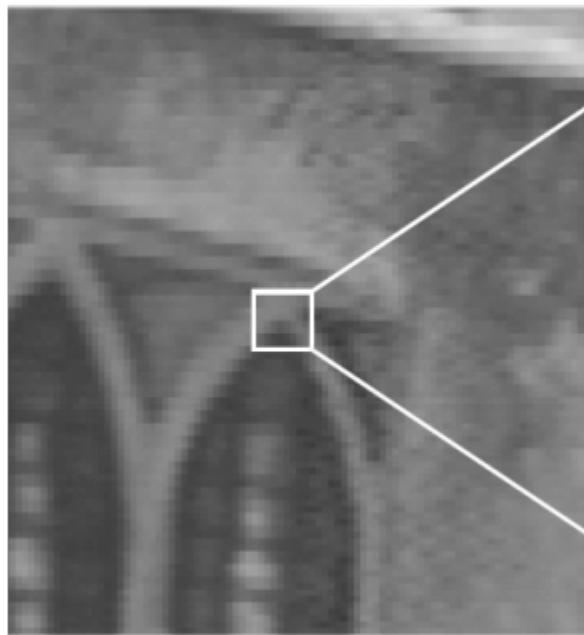
“corner”: significant change
in all directions

Features from Accelerated Segment Test (FAST)

- Compare local neighborhood around each pixel to determine if the pixel is a good feature point
- For each pixel x
 - Look at the pixels on the border of a circle of radius r around x
 - Let n be the number of contiguous pixels (check for wrap-around!) whose intensities are either
 - $\text{all } > (I(x) + T)$ or $\text{all } < (I(x) - T)$
 - $I(x)$ is the intensity at pixel x and T is a threshold
 - If $n \geq n^*$ then the pixel is considered a feature
 - Original paper suggest using $r = 3$ (yields 16 pixels on border) and $n^* = 12$
 - Later results suggest using $r = 3$ (yields 16 pixels on border) and $n^* = 9$

FAST Example

- $R=3$ and $n^*=12$



- Techniques can be employed to decrease computation time

FAST Example ($r = 3, n^* = 12$)

Input



T=10



T=30



T=50



FAST Example ($r = 3, n^* = 9$)

Input



T=10



T=30



T=50



Harris Detector

- Based on the matrix M , which is a 2×2 correlation matrix from image derivatives around a pixel (recall good feature of KLT)

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Sum over image region-
the area we are checking for corner

Window/Weighting function

Gradients with respect to x and y

- To evaluate the “cornerness” of a particular pixel location from its local region

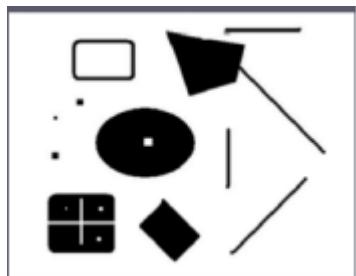
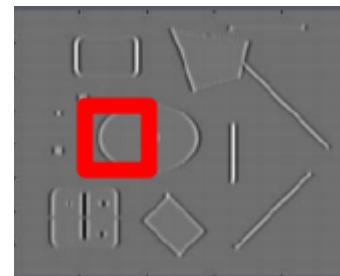
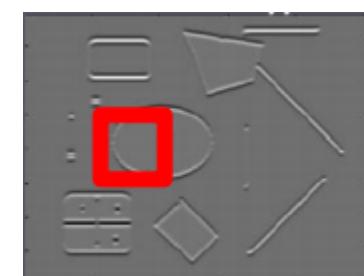


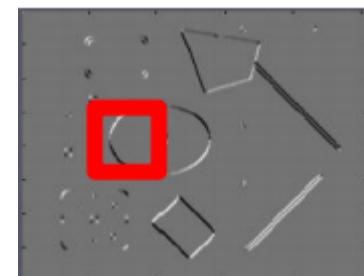
Image I



I_x



I_y



$I_x I_y$

Gradients

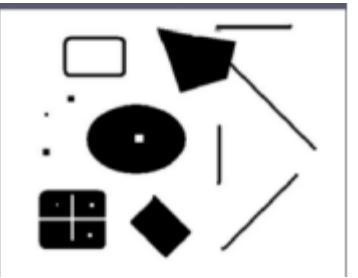
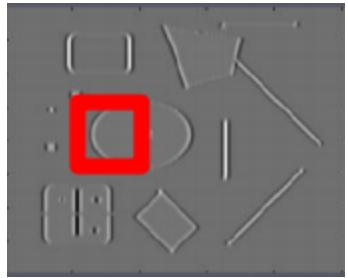
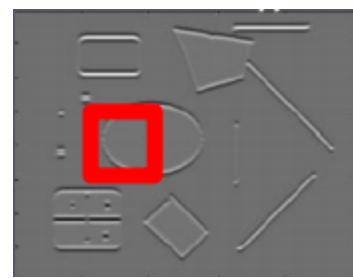


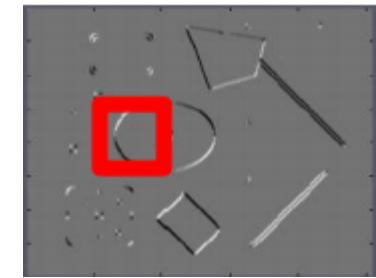
Image I



I_x



I_y



$I_x I_y$

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

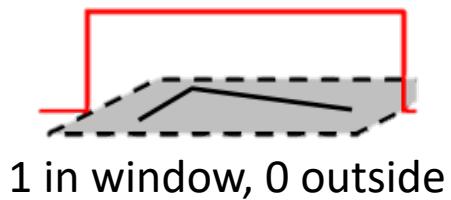
Window/Weighting Function $w(x, y)$

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Option 1: uniform window

- Sum over square window

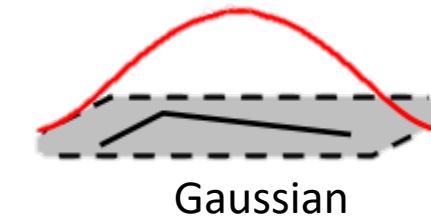
- $M = \sum_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$



- Option 2: Gaussian window

- Gaussian performs weighted sum

- $M = \sum_{x,y} g(\sigma) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$

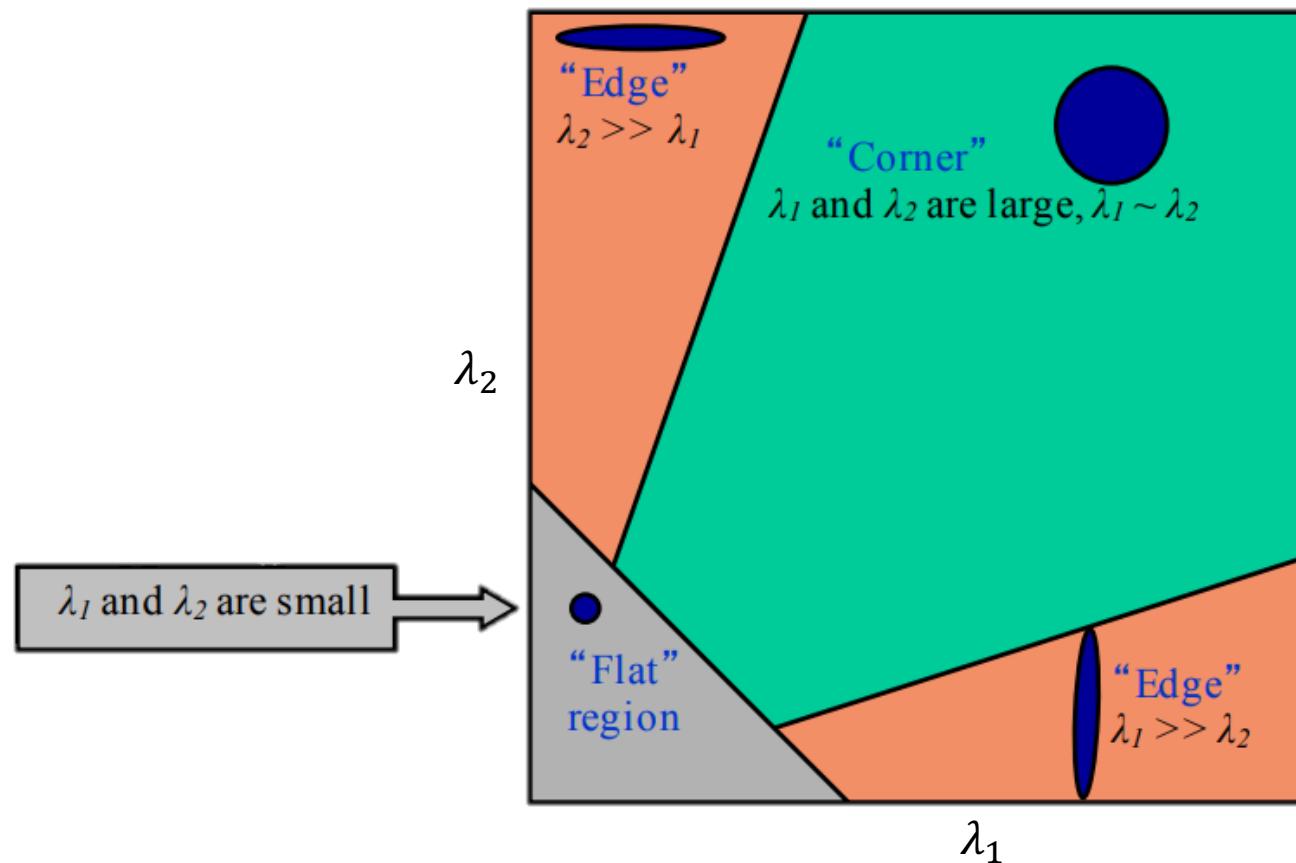


Eigenvalues of M

- Let $\lambda_1 \geq \lambda_2$ be the Eigenvalues of M
 - $\lambda_1 < e$: intensity is nearly constant over patch
 - $\lambda_1 \gg \lambda_2$: edge was found
 - $\lambda_2 > T$: corner or textured patch found

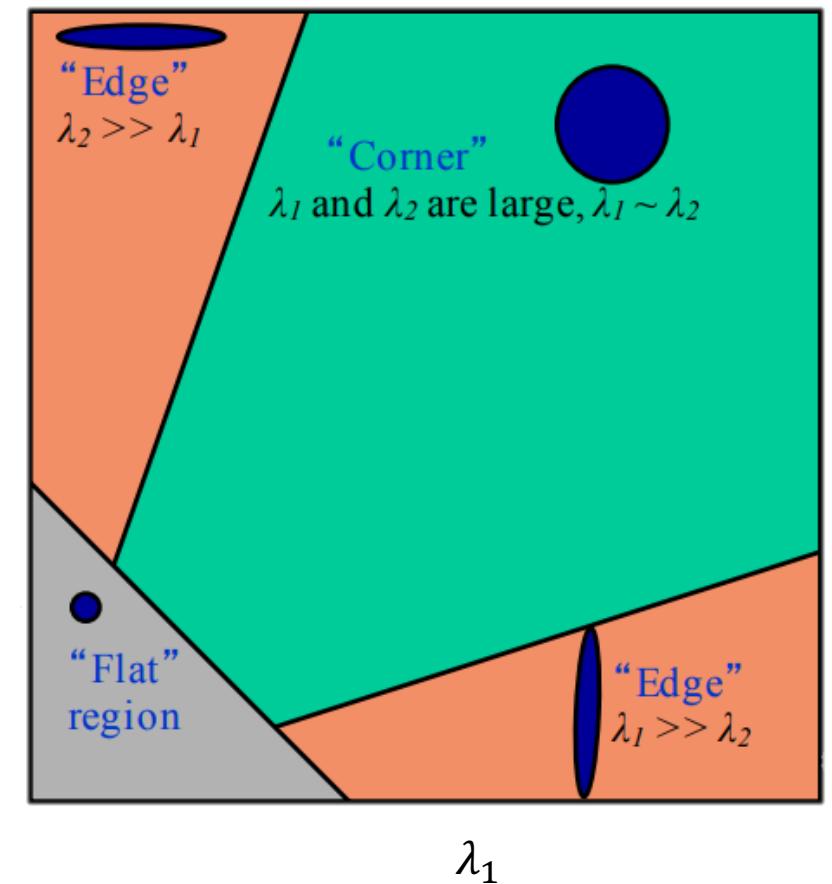
Interpreting the Eigenvalues

- Classification of image points using eigenvalues of M
 - No ordering of eigenvalues



Harris Corner Response Function

- $R = \det(M) - \alpha \text{trace}(M)^2 = \lambda_1\lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$
 - α :constant (0.04 to 0.06)
- Fast approximation
 - R depends on only eigenvalue of M
 - But, do not compute them (still fast in the 80's)



Harris Algorithm

1. Compute Gaussian derivative at each pixel
2. Compute second moment matrix M in a Gaussian window around each pixel
3. Compute corner response function R
4. Threshold R
5. Find local maxima of response function (nonmaximum suppression)

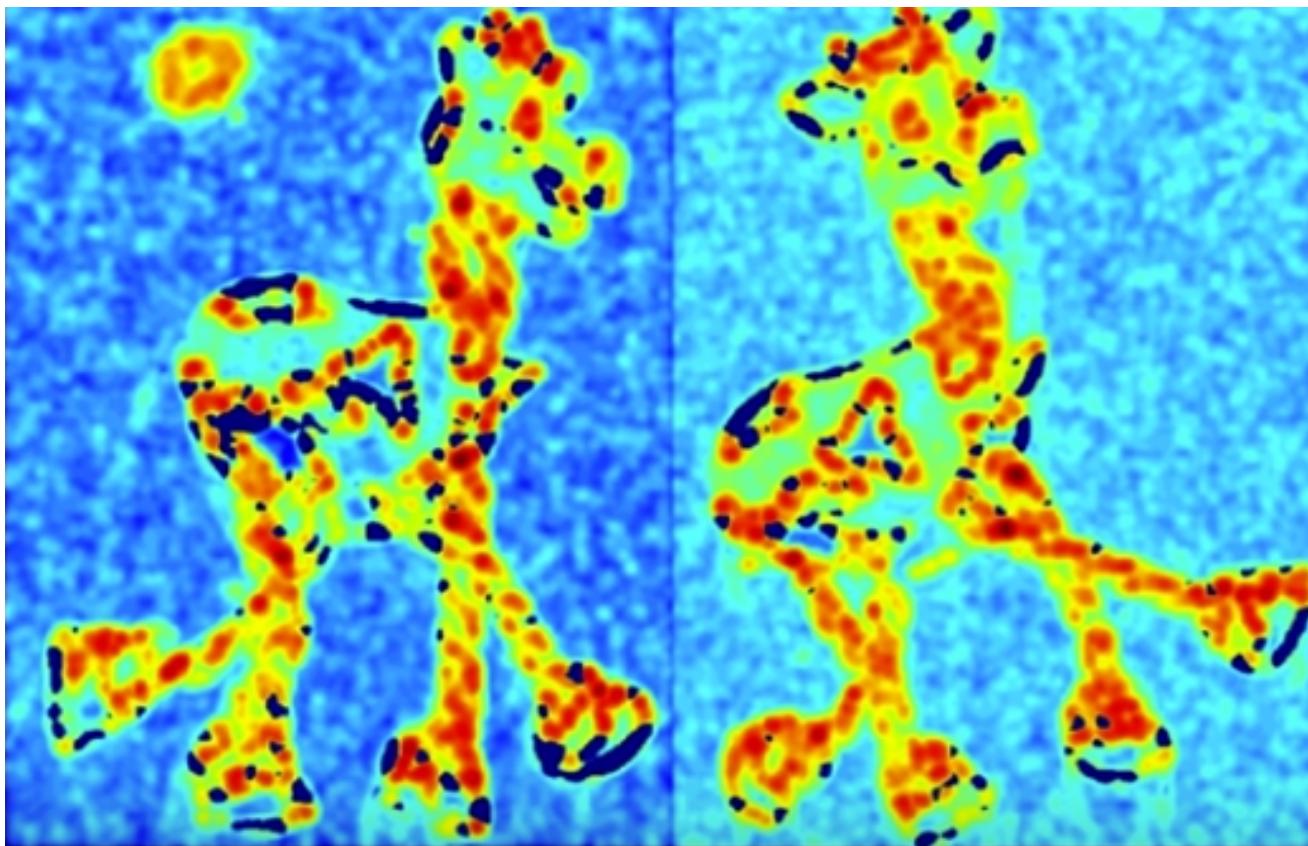
Harris Detector Example

- Head has been rotated



Harris Detector Example

- R function



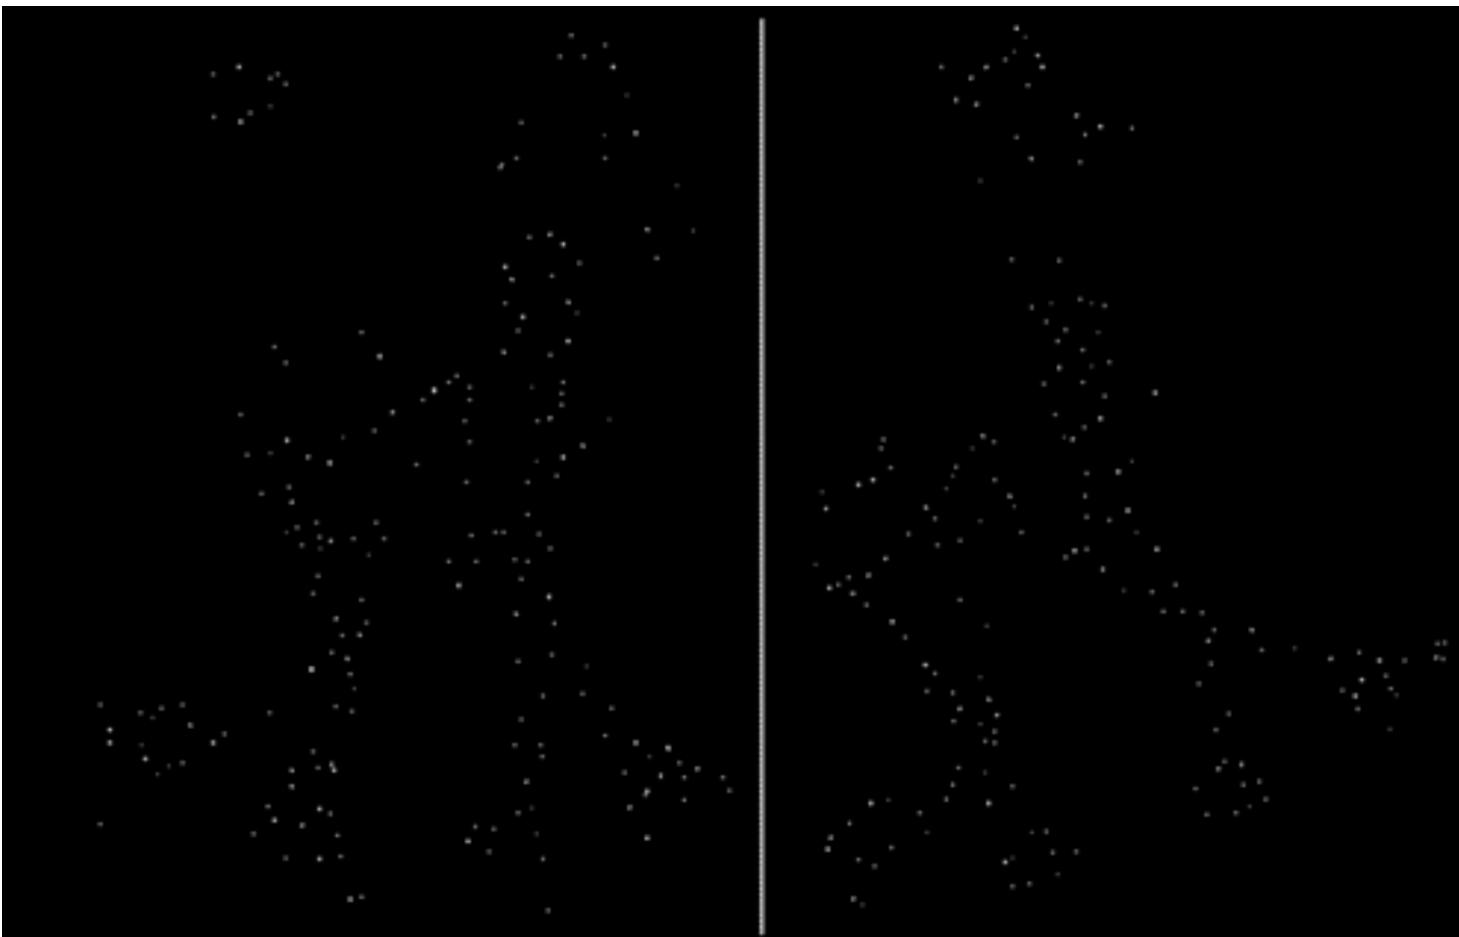
Harris Detector Example

- Thresholding



Harris Detector Example

- Take only the point of local maxima of R



Harris Detector Example

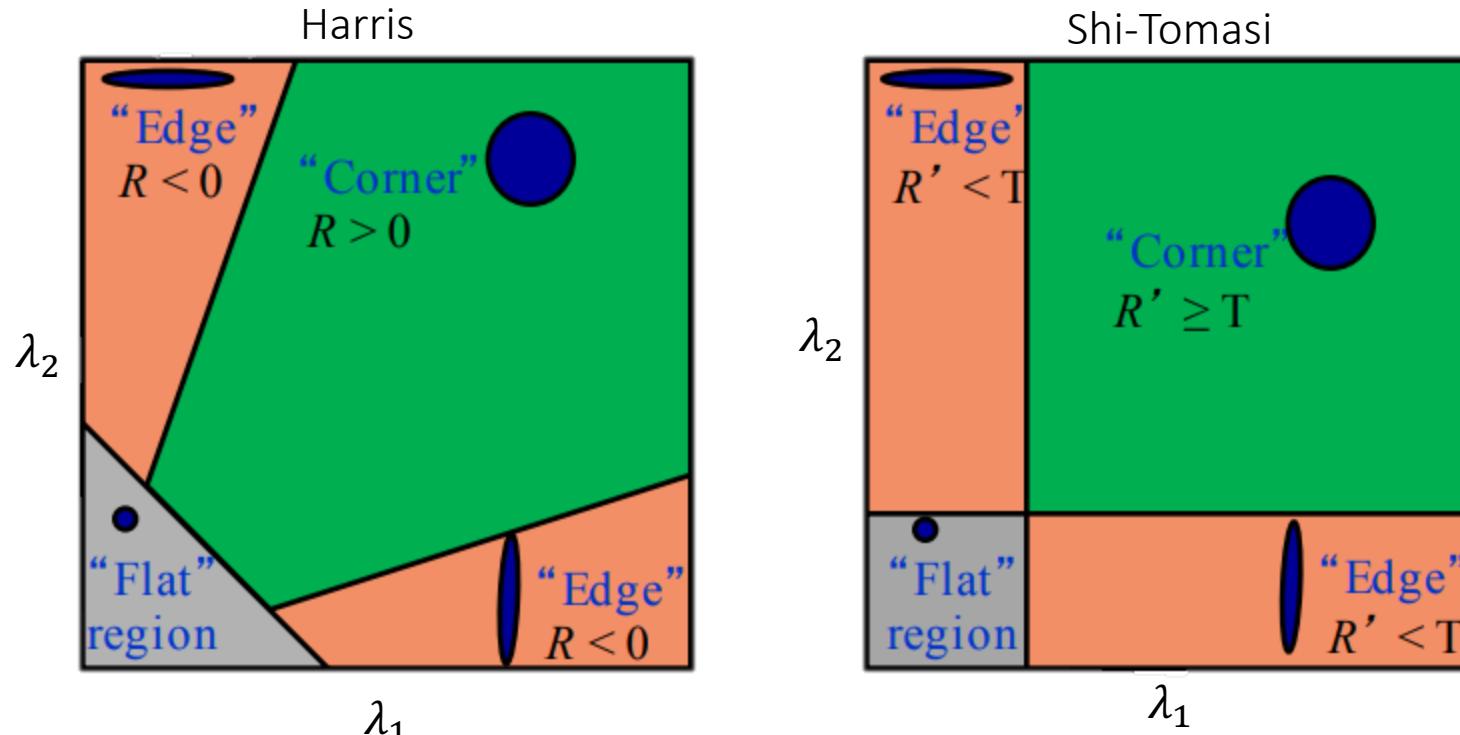


Harris Detector - Responses



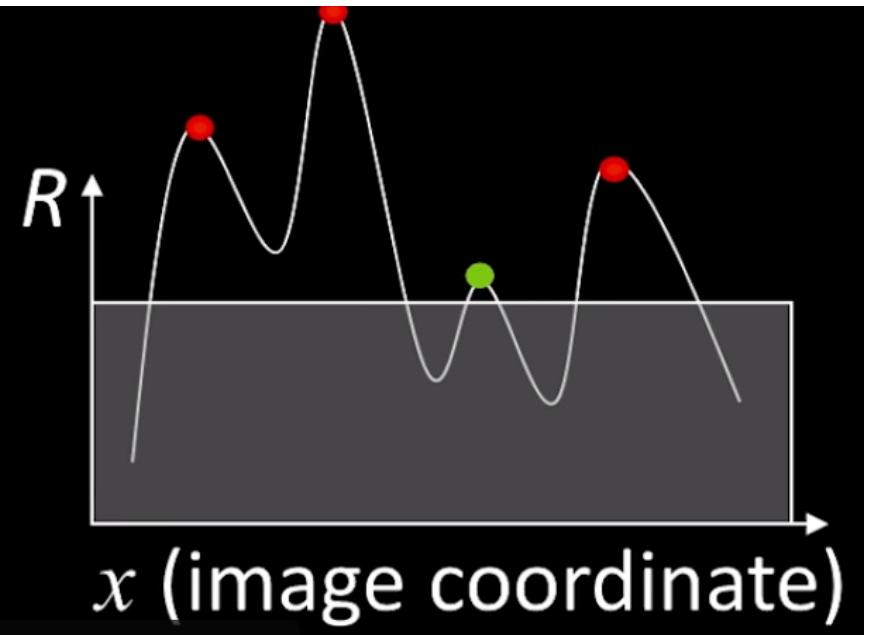
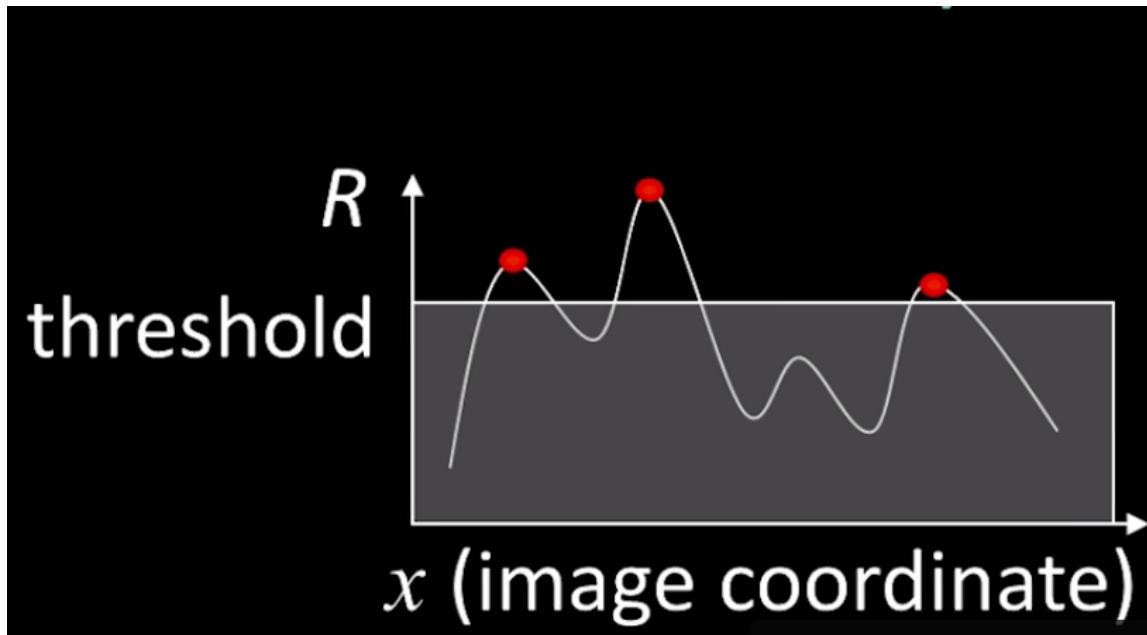
Shi-Tomasi Corner Detector

- Similar to Harris corner points
- Harris: $R = \det(M) - \alpha \text{trace}(M)^2 = \lambda_1\lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$
- Shi-Tomasi (KLT good feature): $R' = \min(\lambda_1, \lambda_2)$



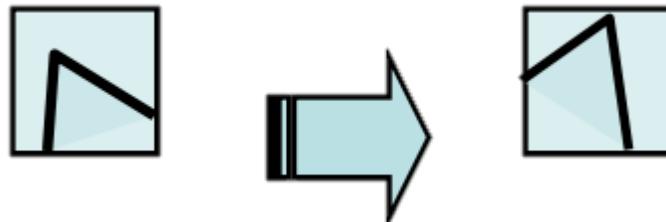
Harris and Shi-Tomasi Detector : Properties

- Intensity invariance?



Harris and Shi-Tomasi Detector : properties

- Intensity invariance
- Rotation invariance?

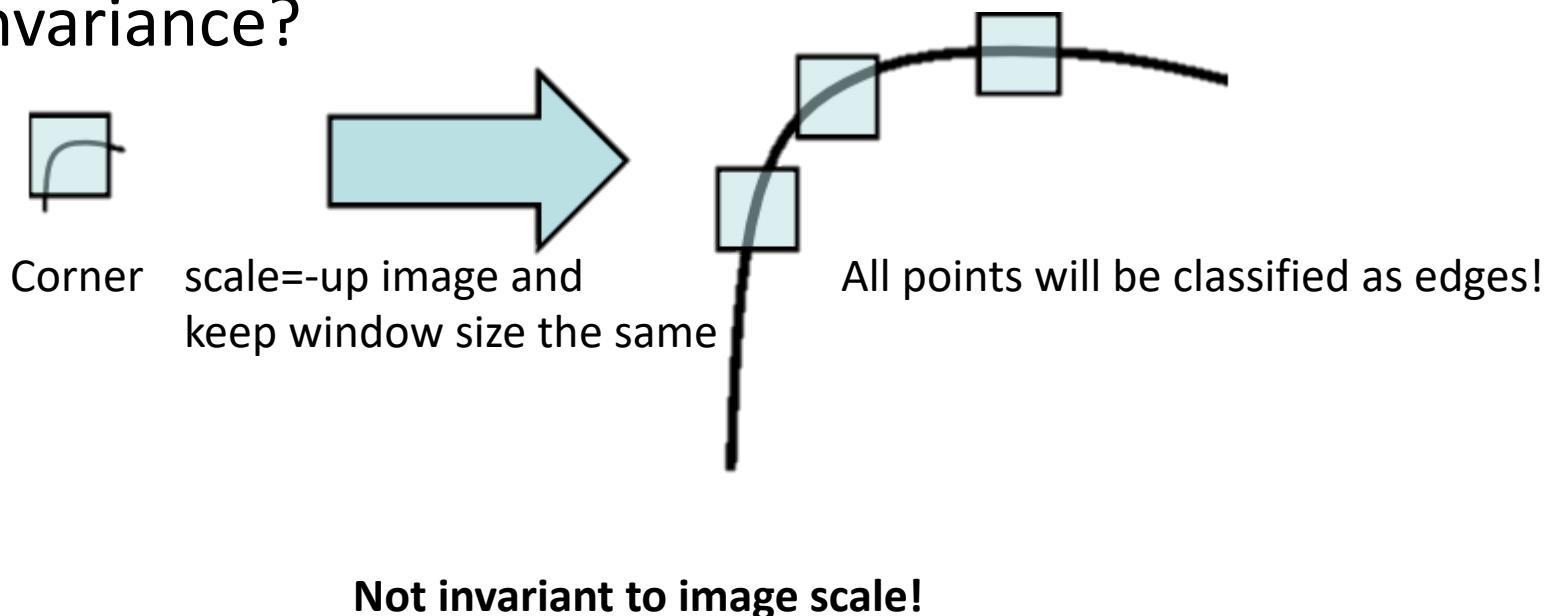


Eigenvalues remain the same

Corner response R is invariant to image rotation

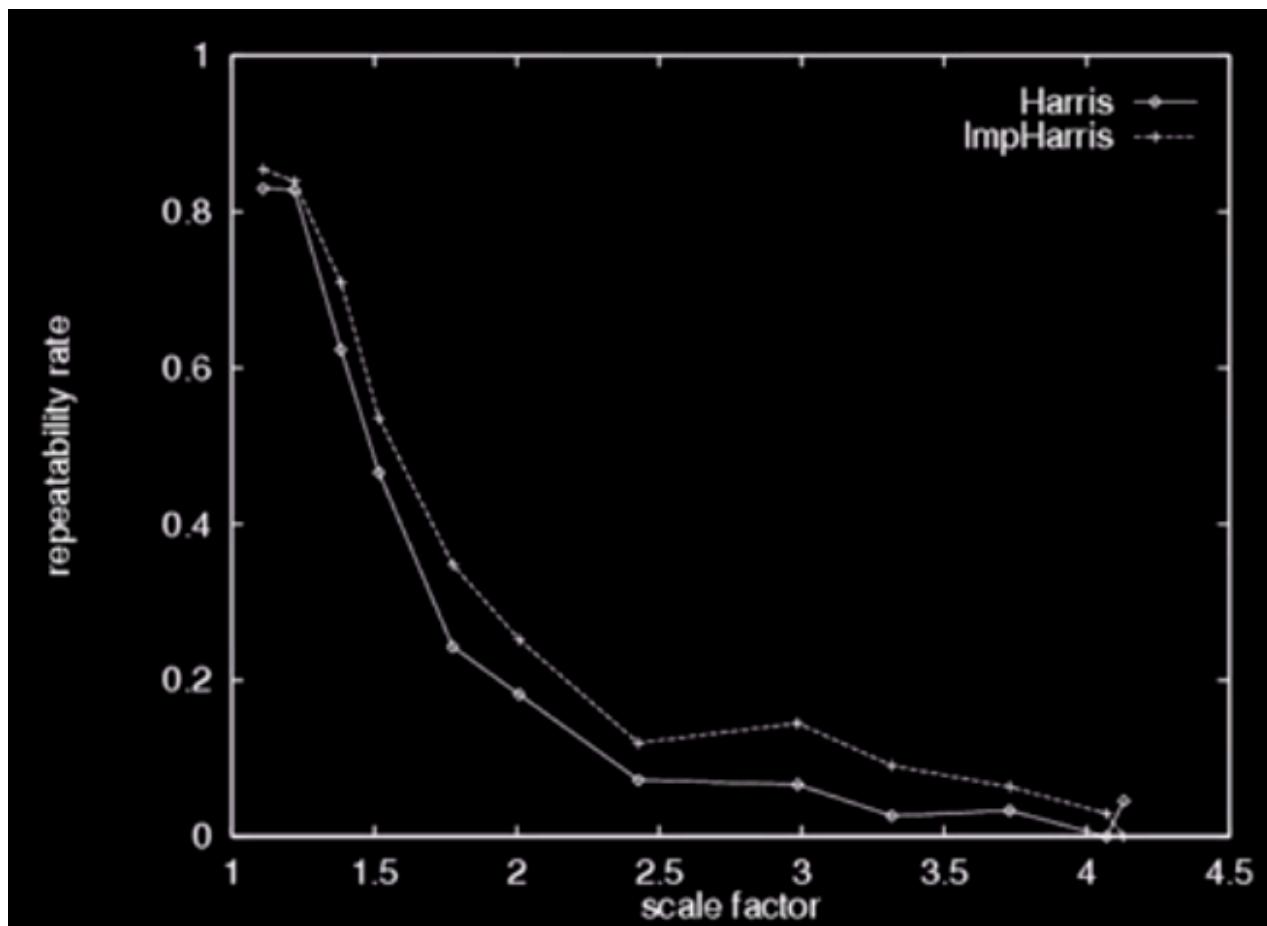
Harris and Shi-Tomasi Detector : properties

- Intensity invariance
- Rotation invariance
- Scale invariance?



Harris and Shi-Tomasi Detector : Properties

- Repeatability rate v.s. scale factor



Key Takeaways

- Harris corner detection is similar to Shi-Tomasi good feature detector
 - Theoretically, use eigen values to detect the corners (a little bit different equation)
- They are intensity invariant, rotation invariant, but NOT scale invariant

Outline

- Local Invariant Features
 - Motivation
 - Requirements, Invariances
- Keypoint Localization
 - Feature from Accelerated Segment Test (FAST)
 - Harris
 - Shi-Tomasi
- **Scale Invariant Region Selection**
 - Automatic scale selection
 - Laplacian-of-Gaussian detector
 - Difference-of-Gaussians detector
- Local Descriptors
 - Orientation normalization
 - SIFT

From Points to Regions

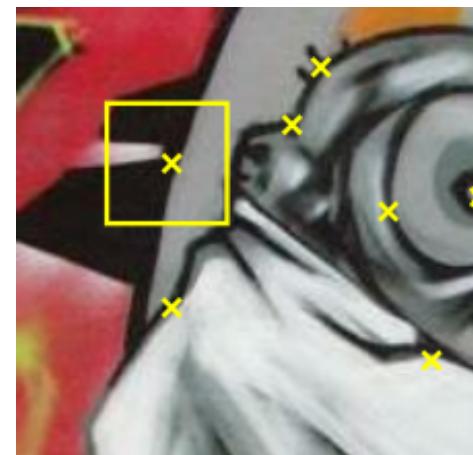
- Harris and Shi-Tomasi operators define interest points



- How can we detect scale invariant interest region?
 - In order to compare points, we need to compute a descriptor over a region
 - How can we define such a region in a scale invariant manner?

Naïve Approach: Exhaustive Search

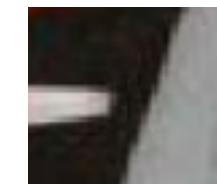
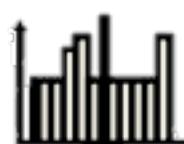
- Multi-scale procedure
 - Compare descriptors while varying the patch size in the other image



f_A similarity measure f_B



\neq



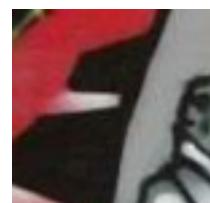
$$d(f_A, f_B)$$

Naïve Approach: Exhaustive Search

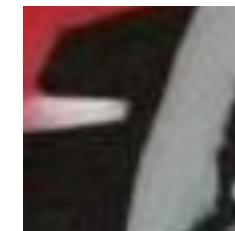
- Multi-scale procedure
 - Compare descriptors while varying the patch size in the other image



f_A similarity measure f_B



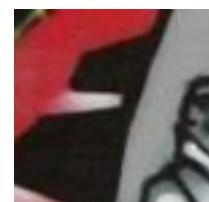
\neq



$$d(f_A, f_B)$$

Naïve Approach: Exhaustive Search

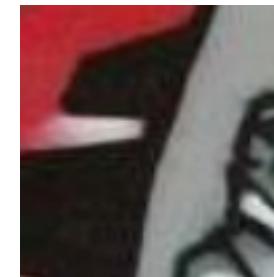
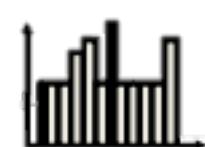
- Multi-scale procedure
 - Compare descriptors while varying the patch size in the other image



f_A similarity measure f_B

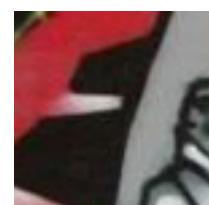
\neq

$$d(f_A, f_B)$$



Naïve Approach: Exhaustive Search

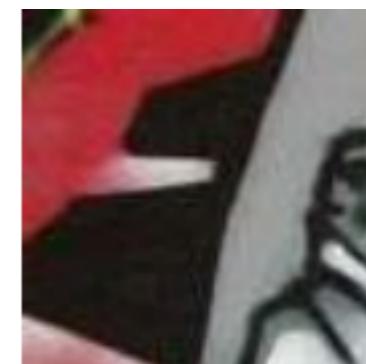
- Multi-scale procedure
 - Compare descriptors while varying the patch size in the other image



=



$$d(f_A, f_B)$$



Naïve Approach: Exhaustive Search

- Comparing descriptors while varying the patch size
 - Computationally inefficient
 - But possible for matching
 - Prohibitive for retrieval in large databases



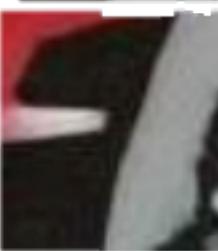
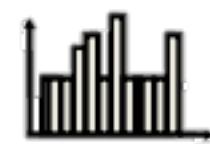
f_A

similarity measure



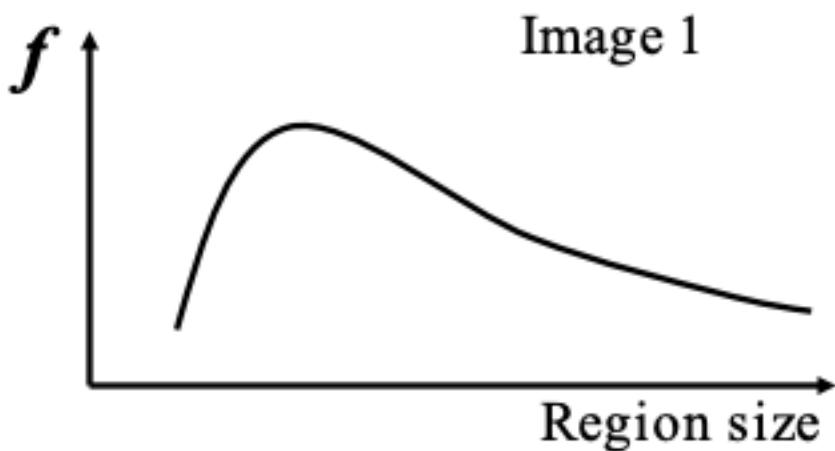
=

$d(f_A, f_B)$



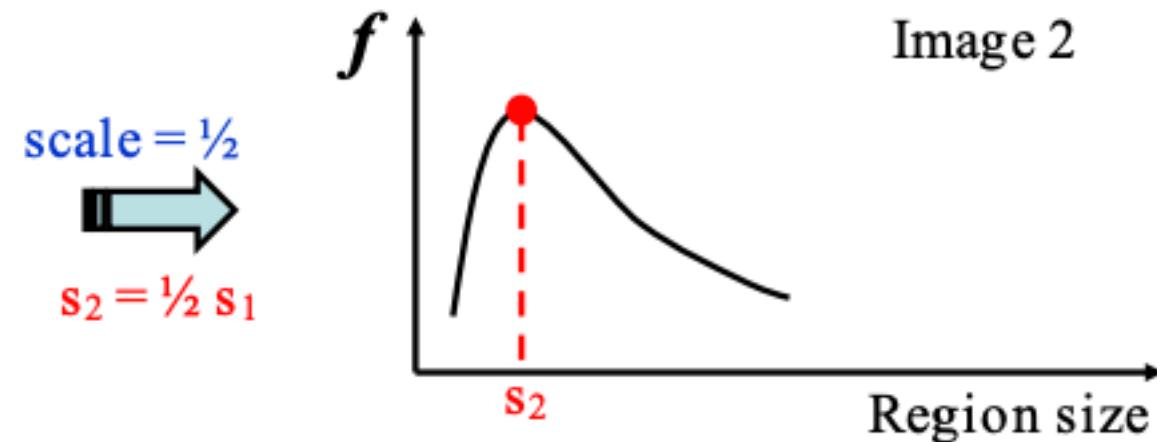
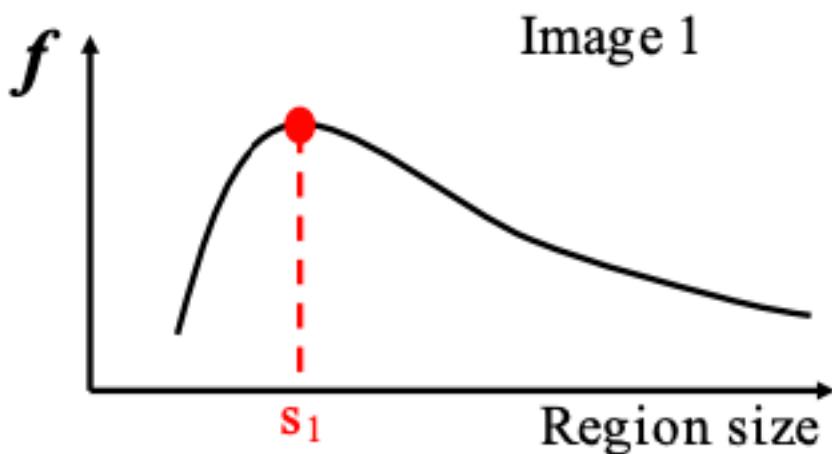
Automatic Scale Selection

- Solution:
 - Design a “magic” function f that is maximal at the correct region size for a given location (indicates the best region size)



Automatic Scale Selection

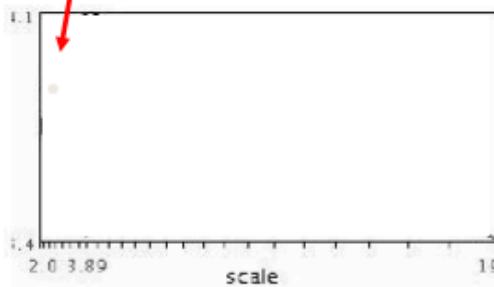
- Choose the region size that is best
 - Important: this scale invariant region size (for a point) is found in each image independently!



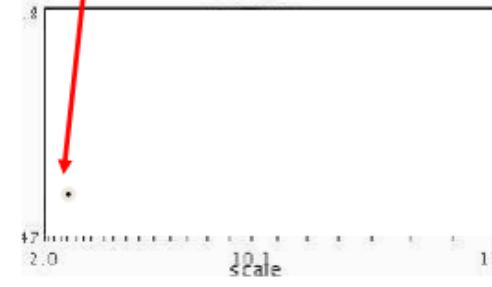
$$\begin{array}{l} \text{scale} = \frac{1}{2} \\ \Rightarrow \\ s_2 = \frac{1}{2} s_1 \end{array}$$

Automatic Scale Selection

- Function responses for increasing scale (scale signature)



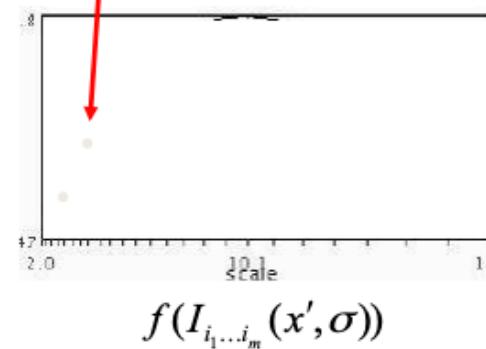
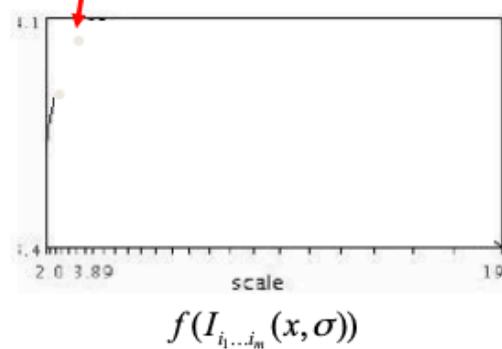
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



$$f(I_{i_1 \dots i_m}(x', \sigma))$$

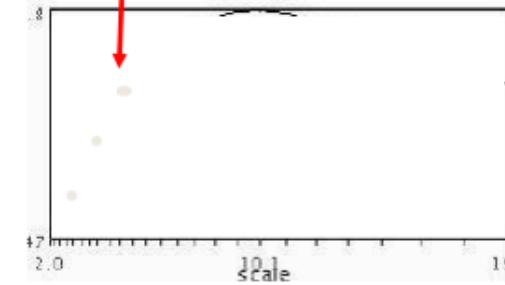
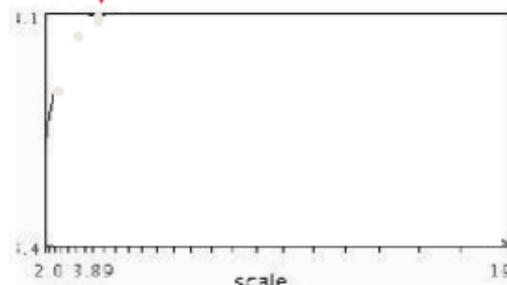
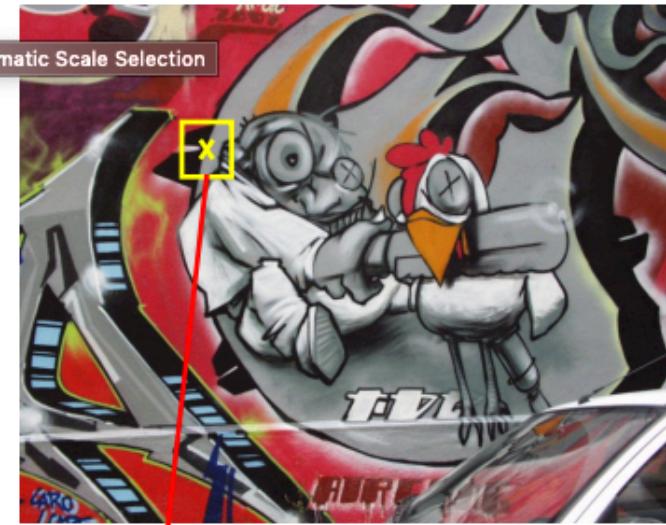
Automatic Scale Selection

- Function responses for increasing scale (scale signature)



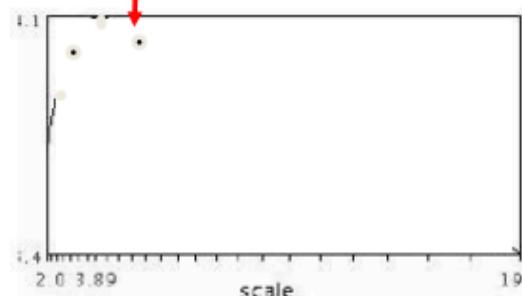
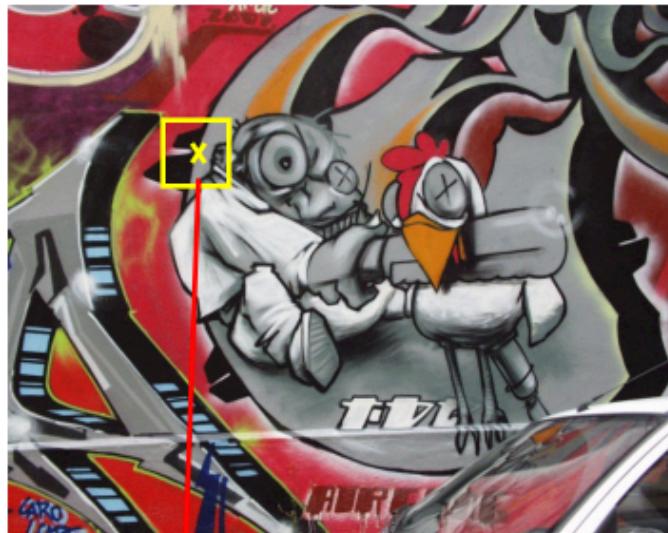
Automatic Scale Selection

- Function responses for increasing scale (scale signature)

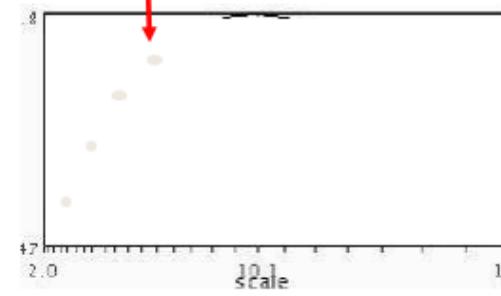


Automatic Scale Selection

- Function responses for increasing scale (scale signature)



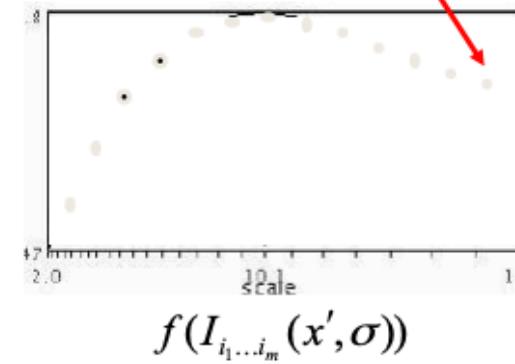
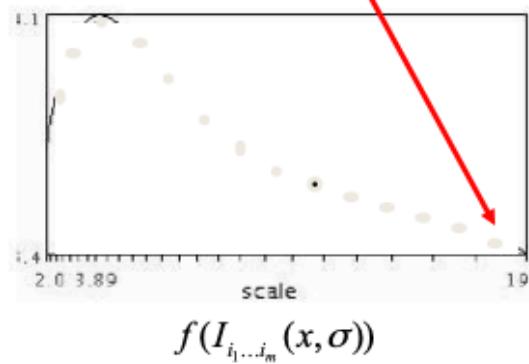
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



$$f(I_{i_1 \dots i_m}(x', \sigma))$$

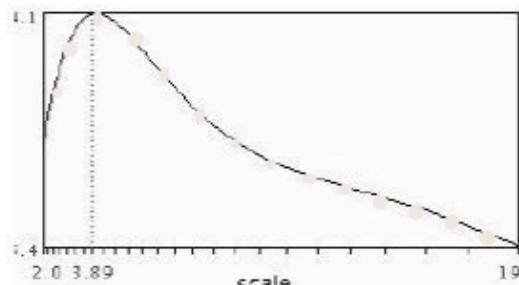
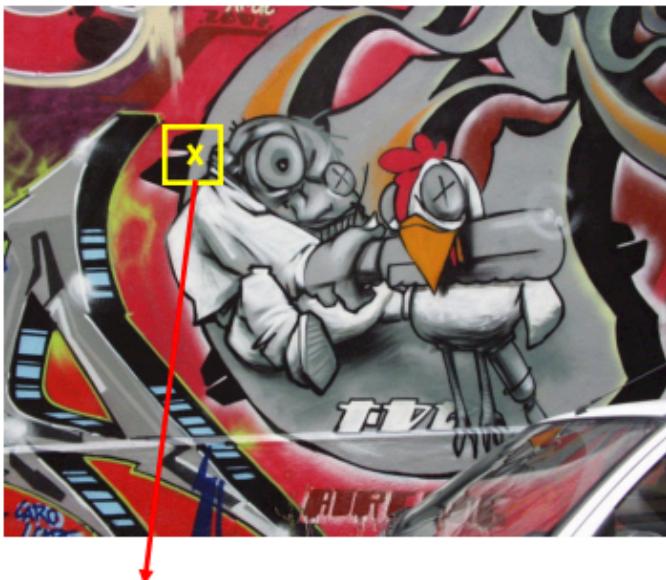
Automatic Scale Selection

- Function responses for increasing scale (scale signature)

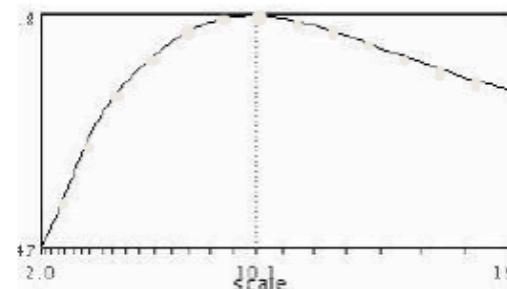


Automatic Scale Selection

- Function responses for increasing scale (scale signature)



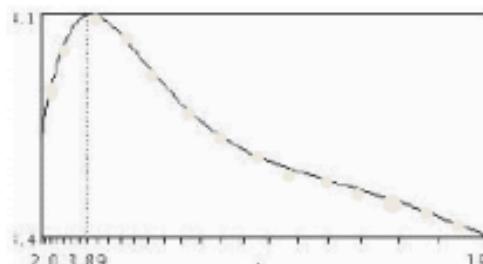
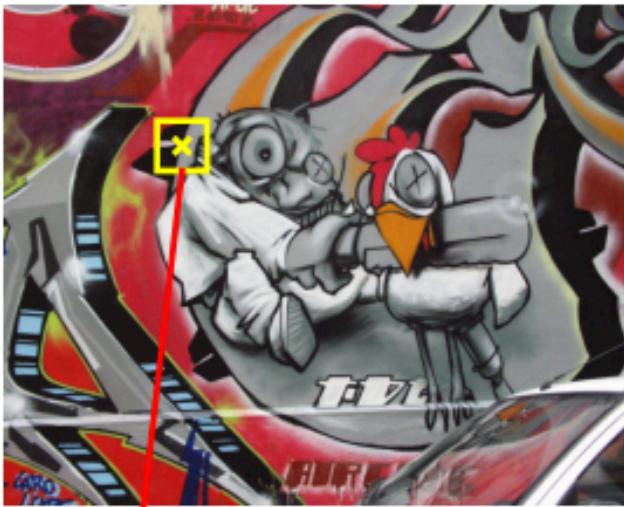
$$f(I_{i_1 \dots i_m}(x, \sigma))$$



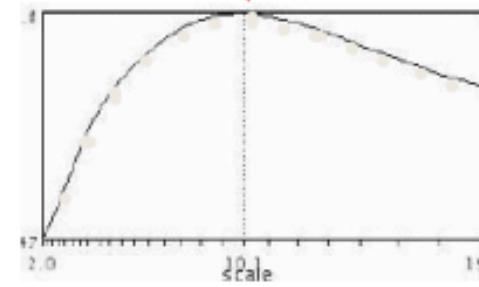
$$f(I_{i_1 \dots i_m}(x', \sigma'))$$

Automatic Scale Selection

- Normalize: Rescale best region size to fixed size



$$f(I_{l_1\dots l_m}(x, \sigma))$$



$$f(I_{l_1\dots l_m}(x', \sigma'))$$

Image1

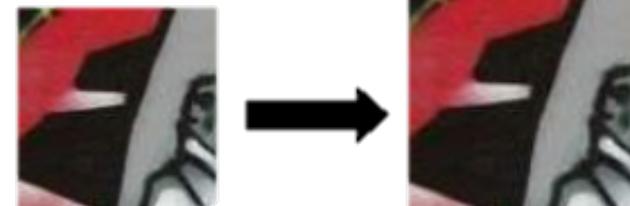
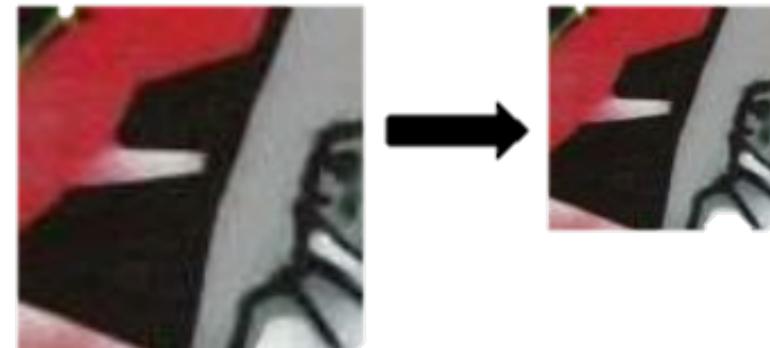


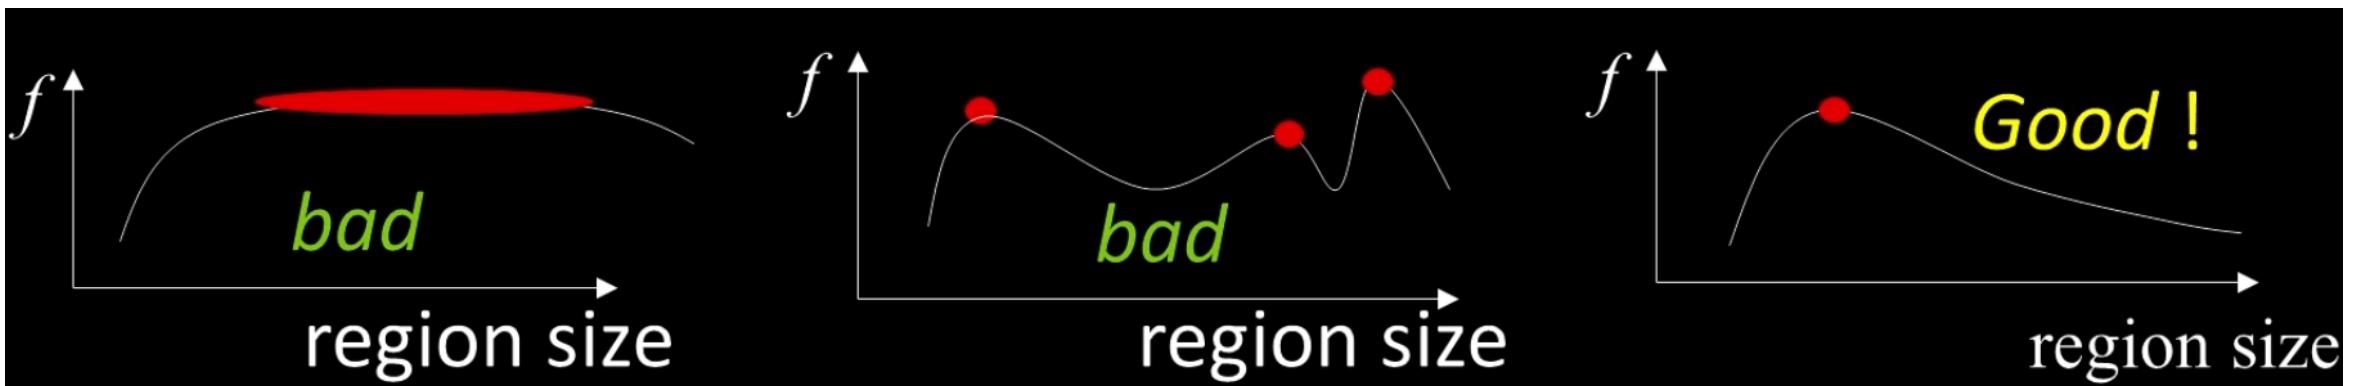
Image2



What could be a possible magic function for f ?

Scale Invariant Detection

- A “good” function for scale detection: has one stable sharp peak



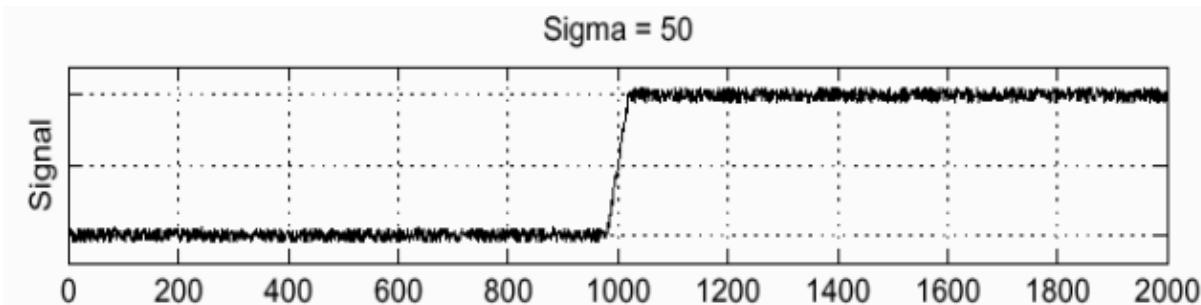
- For usual images: a good function would be a one which responds to contrast (sharp local intensity change)

Blob Detector

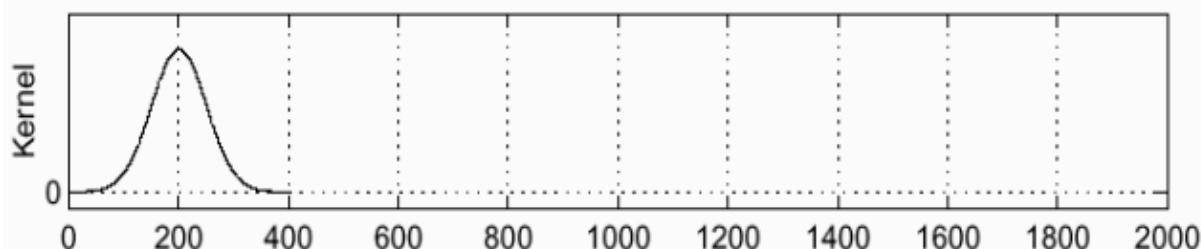


Edge Detector

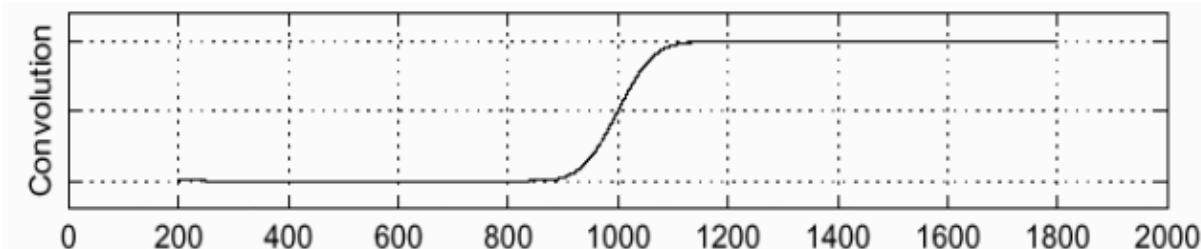
f



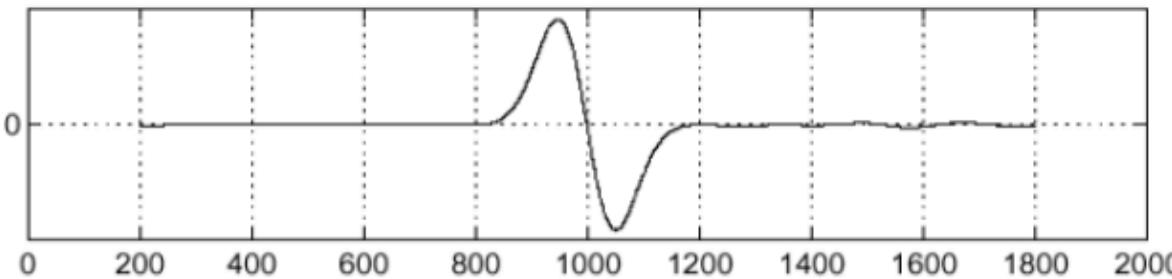
g



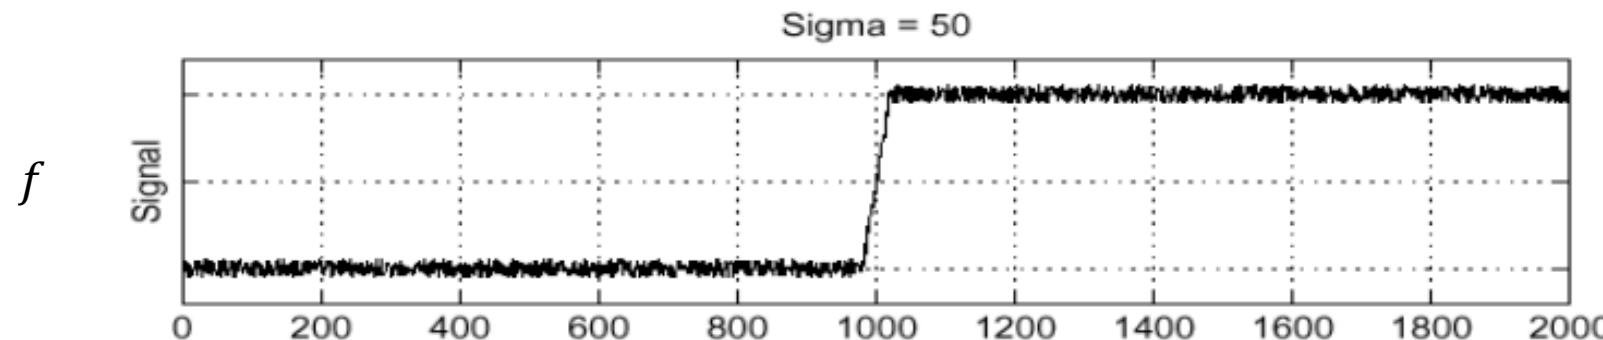
$f * g$



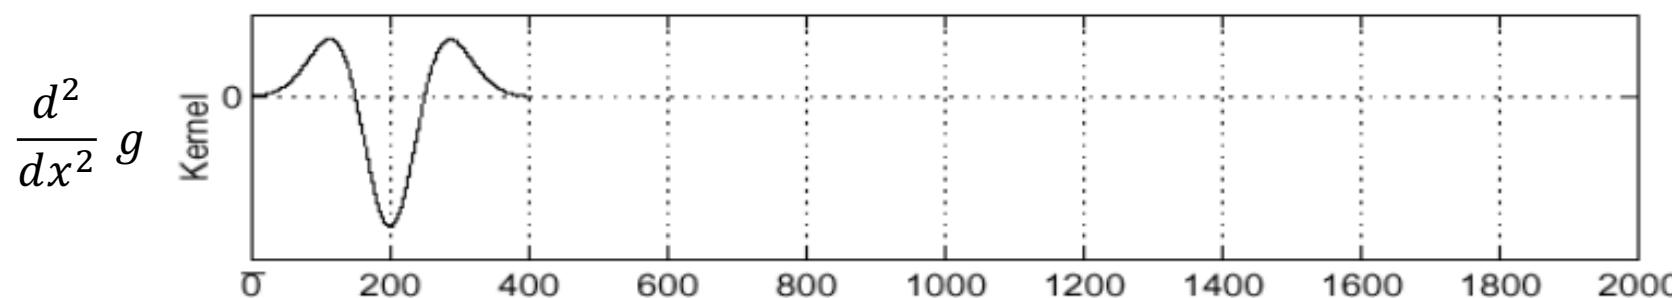
$\frac{d^2}{dx^2} (f * g)$



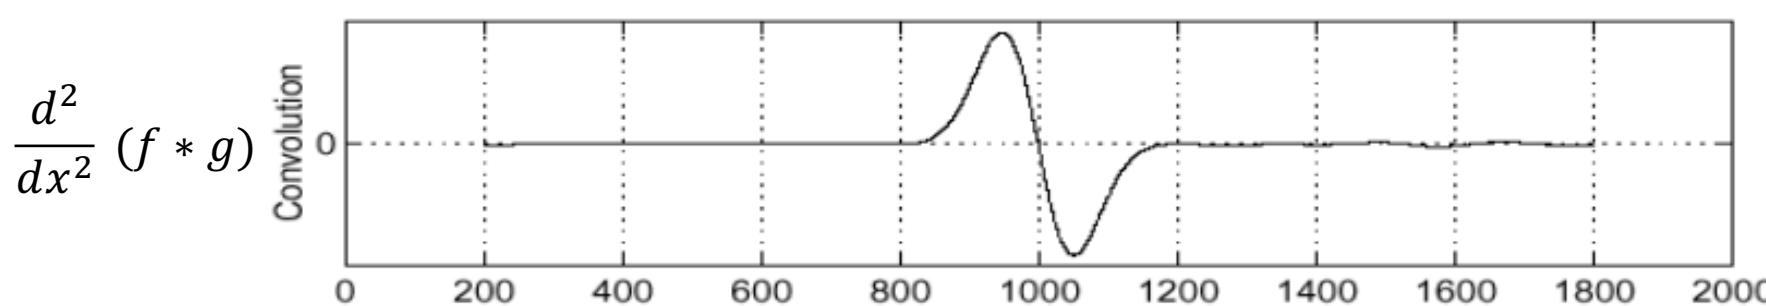
Edge Detector as Zero Crossing



Edge

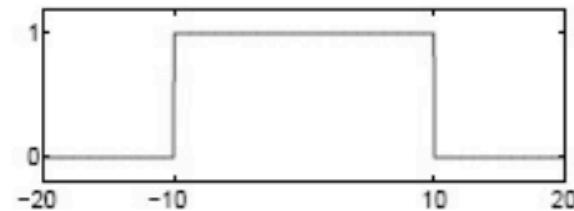


Laplacian

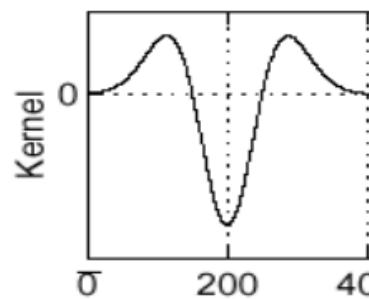


Edge = zero crossing
of second derivative

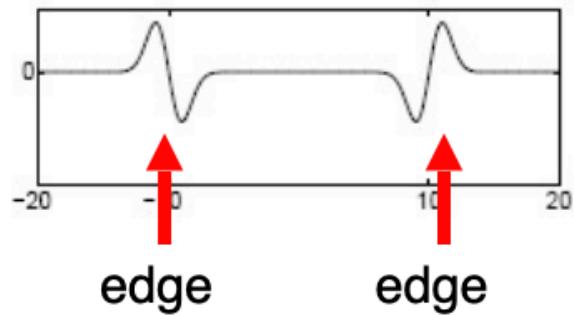
Edge Detector as Zero Crossing



*

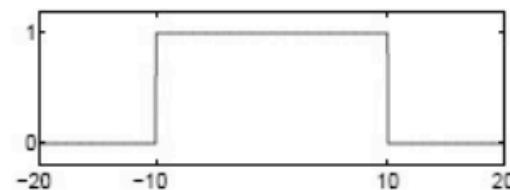


=

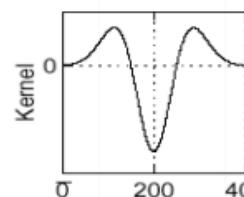


From Edge to Blob

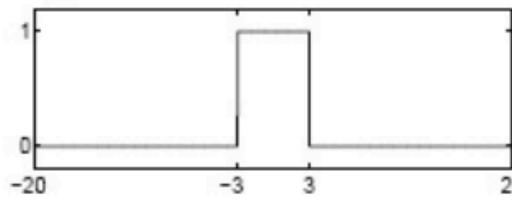
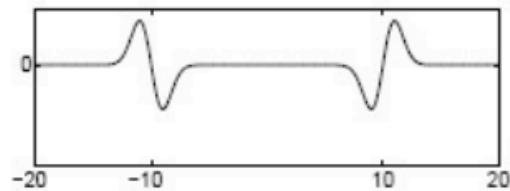
- Can we use the Laplacian to find a blob (rect function in 1D)



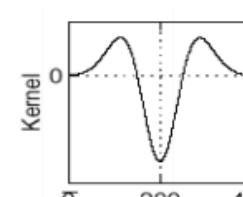
*



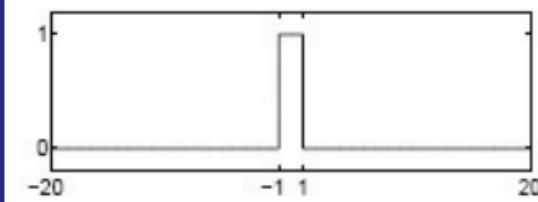
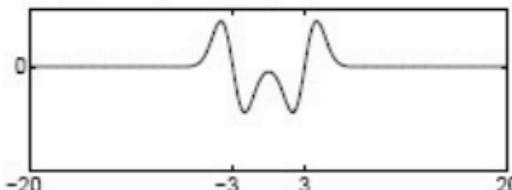
=



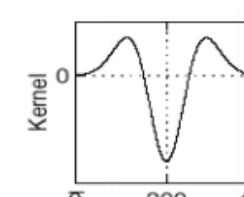
*



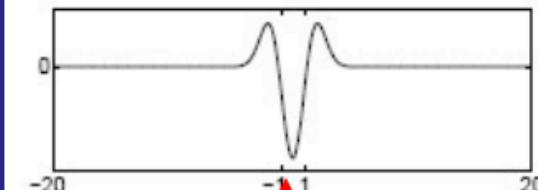
=



*



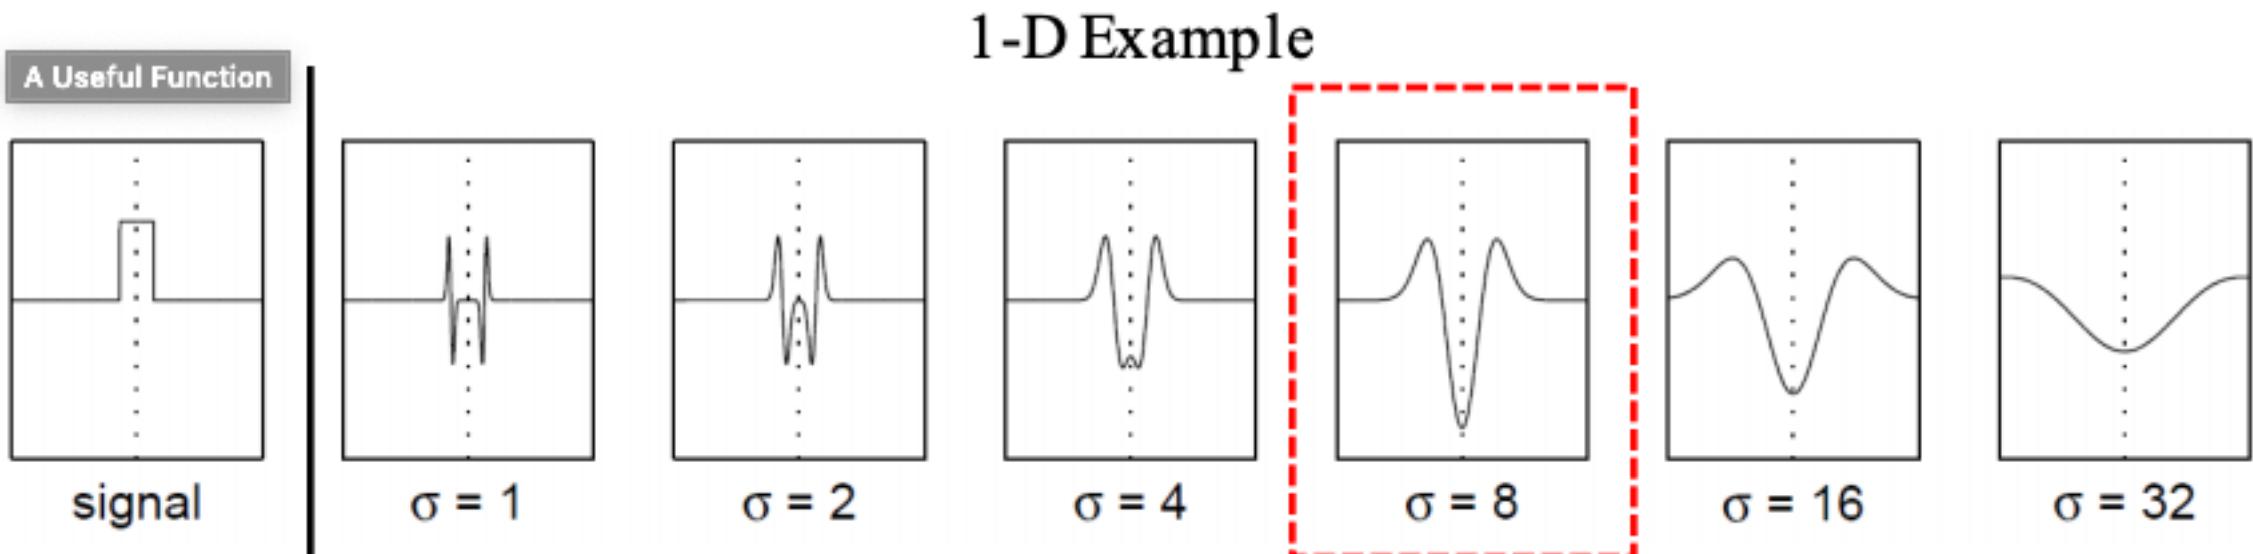
=



maximum

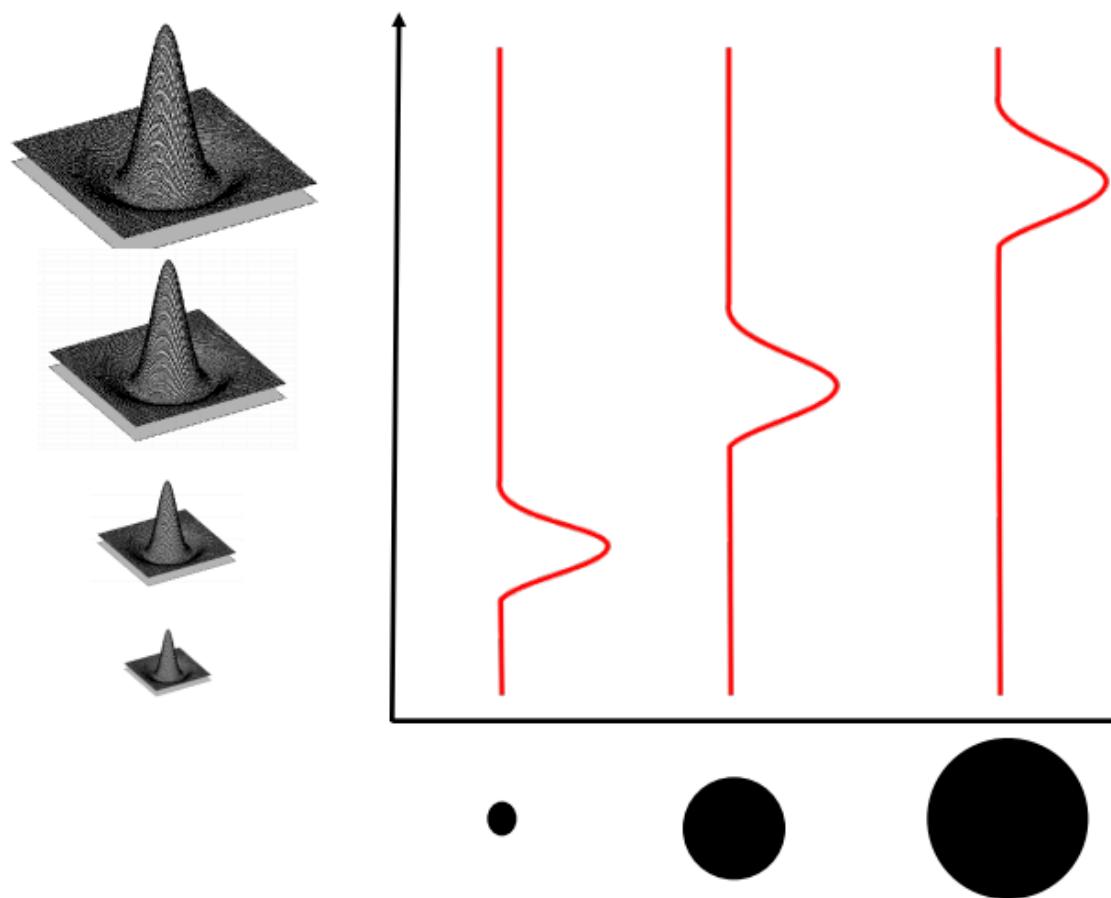
Scale-Normalized Laplacian-of-Gaussian

- $L = \sigma^2(G_{xx}(x, y; \sigma) + G_{yy}(x, y; \sigma))$
 - Scale-normalized Laplacian shows the strongest response at the characteristic scale of the original signal ($\sigma = 8, x = 0$)



A Useful Function

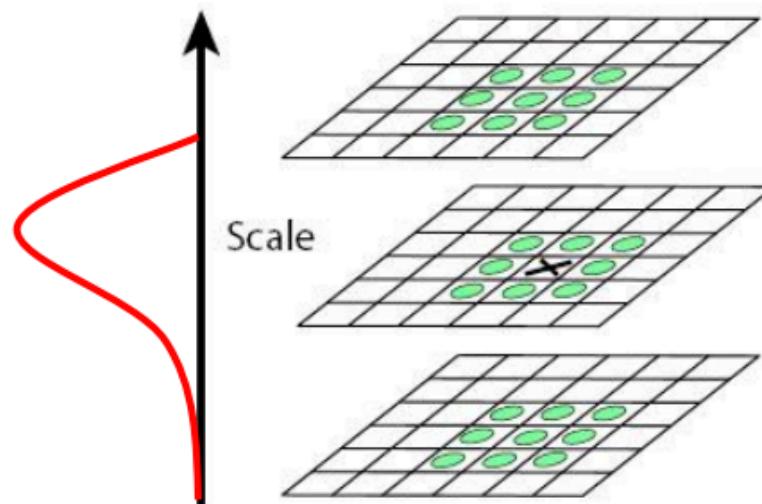
- “Scale-normalized” Laplacian-of-Gaussian = blob/spot detector



Scale Blob Detector

- Convolve image with scale-normalized Laplacian at several scales
- Find maximum of squared Laplacian response in scale-space

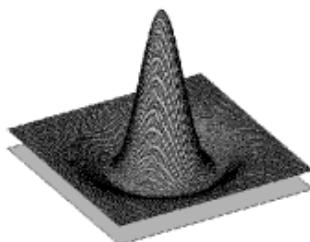
The maxima indicate that a blob has been detected and what's its intrinsic scale



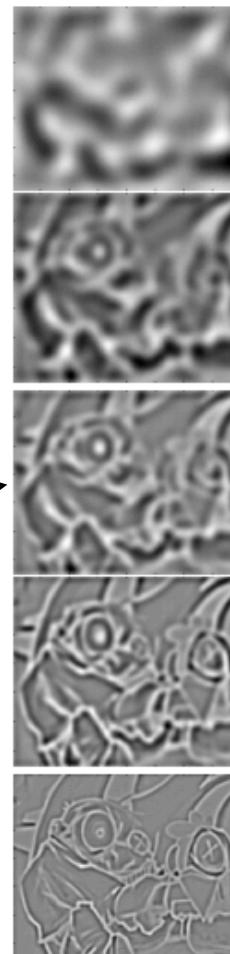
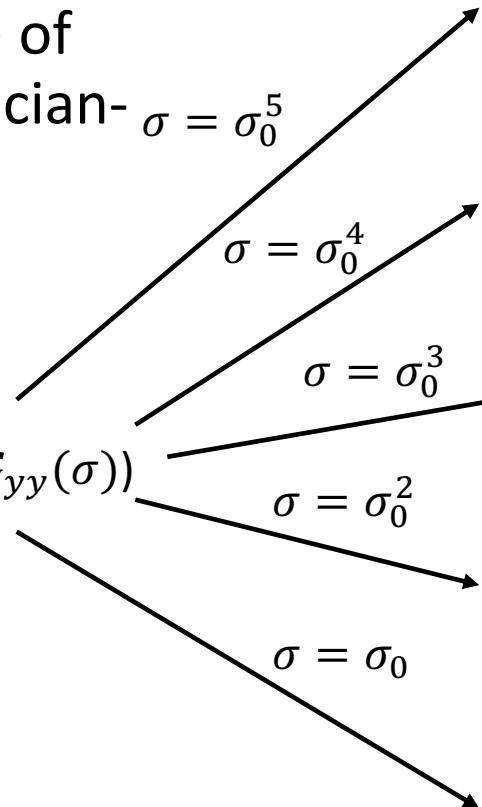
Laplacian-of-Gaussian (LoG)

- Interest points:

- Local maxima in scale space
(above/below/around) of
scale-normalized Laplacian- $\sigma = \sigma_0^5$
of-Gaussian



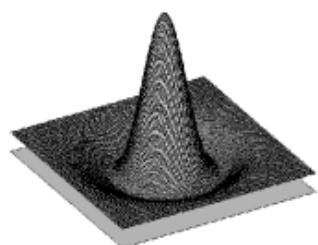
$$\sigma^2(G_{xx}(\sigma) + G_{yy}(\sigma))$$



Laplacian-of-Gaussian (LoG)

- Interest points:

- Local maxima in scale space
(above/below/around) of
scale-normalized Laplacian- $\sigma = \sigma_0^5$
of-Gaussian



$$\sigma^2(G_{xx}(\sigma) + G_{yy}(\sigma))$$

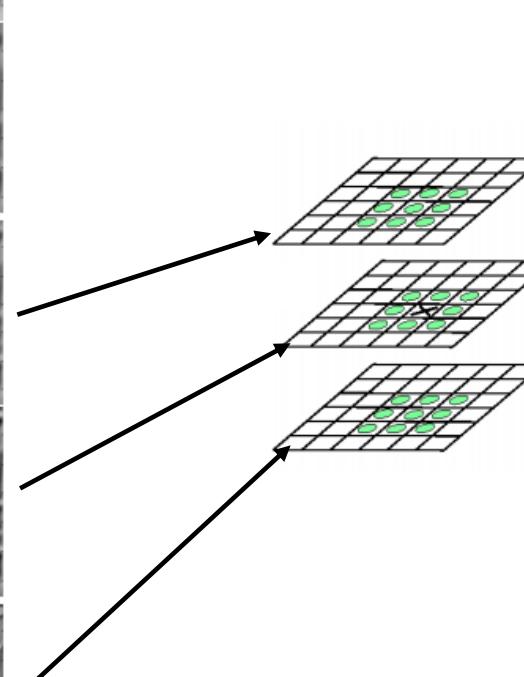
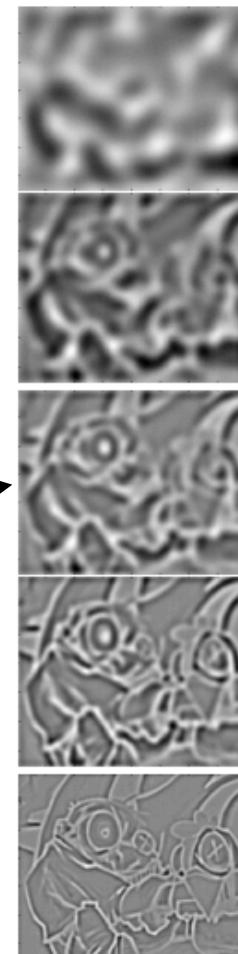
$$\sigma = \sigma_0$$

$$\sigma = \sigma_0^2$$

$$\sigma = \sigma_0^3$$

$$\sigma = \sigma_0^4$$

$$\sigma = \sigma_0^5$$

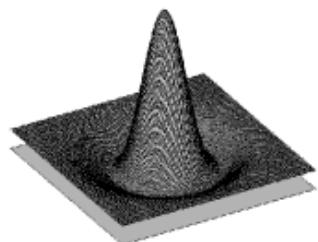


Scale

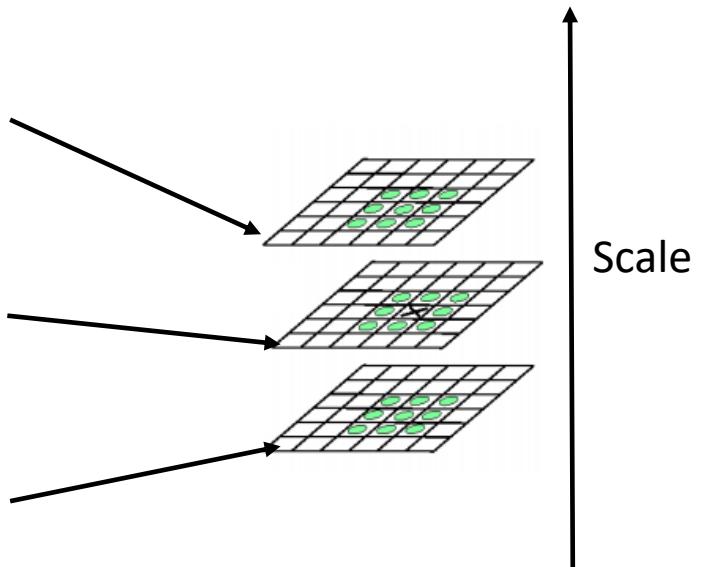
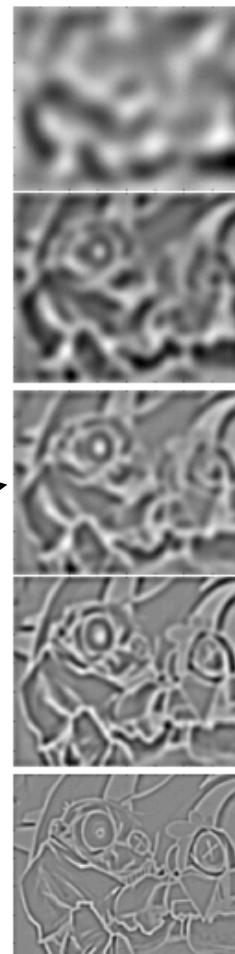
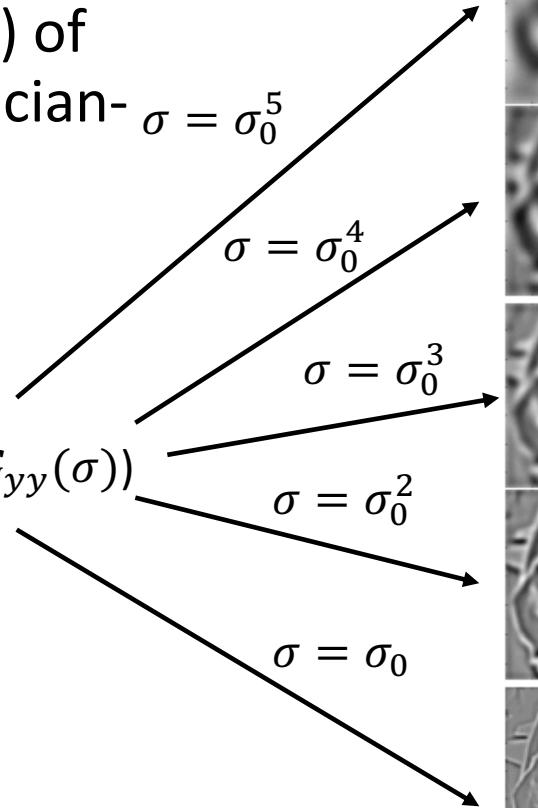
Laplacian-of-Gaussian (LoG)

- Interest points:

- Local maxima in scale space
(above/below/around) of
scale-normalized Laplacian- $\sigma = \sigma_0^5$
of-Gaussian



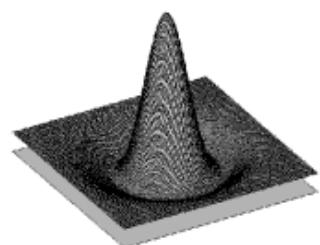
$$\sigma^2(G_{xx}(\sigma) + G_{yy}(\sigma))$$



Laplacian-of-Gaussian (LoG)

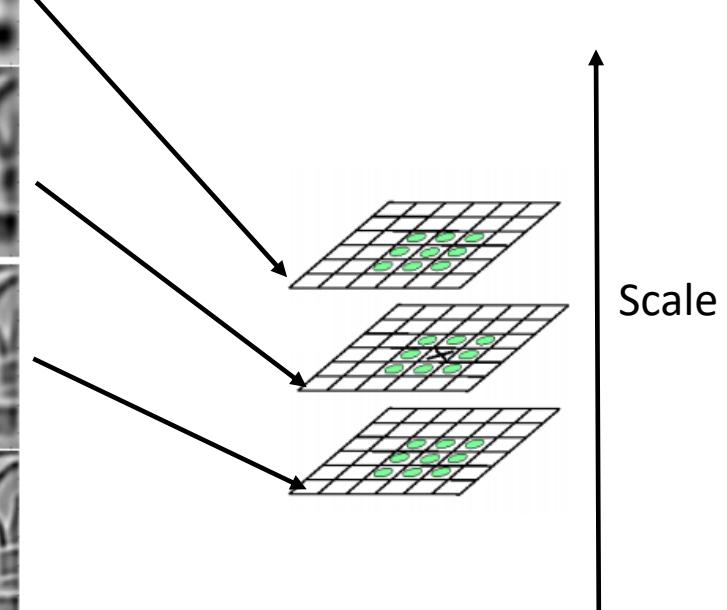
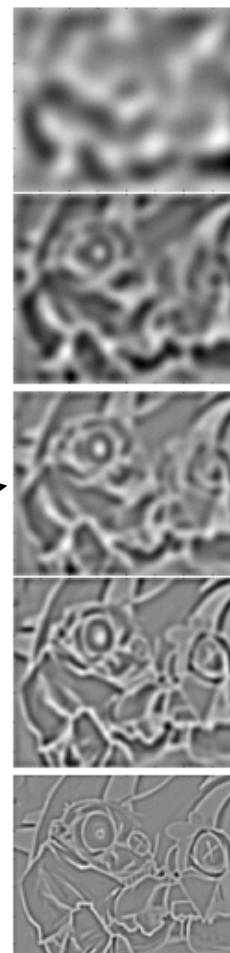
- Interest points:

- Local maxima in scale space
(above/below/around) of
scale-normalized Laplacian- $\sigma = \sigma_0^5$
of-Gaussian



$$\sigma^2(G_{xx}(\sigma) + G_{yy}(\sigma))$$

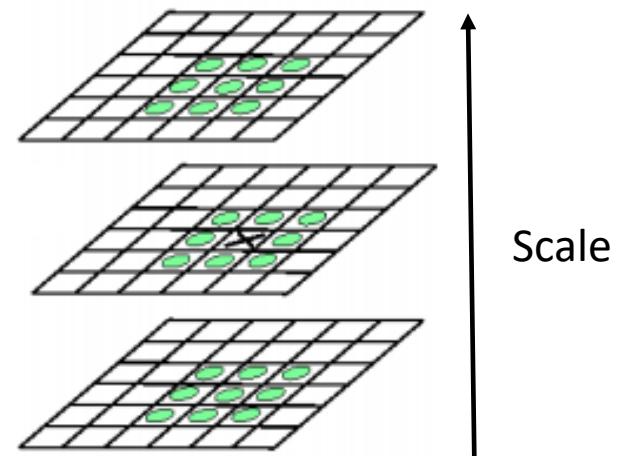
$$\begin{aligned}\sigma &= \sigma_0^5 \\ \sigma &= \sigma_0^4 \\ \sigma &= \sigma_0^3 \\ \sigma &= \sigma_0^2 \\ \sigma &= \sigma_0\end{aligned}$$



Save list of (x, y, σ) that are
locally maximal

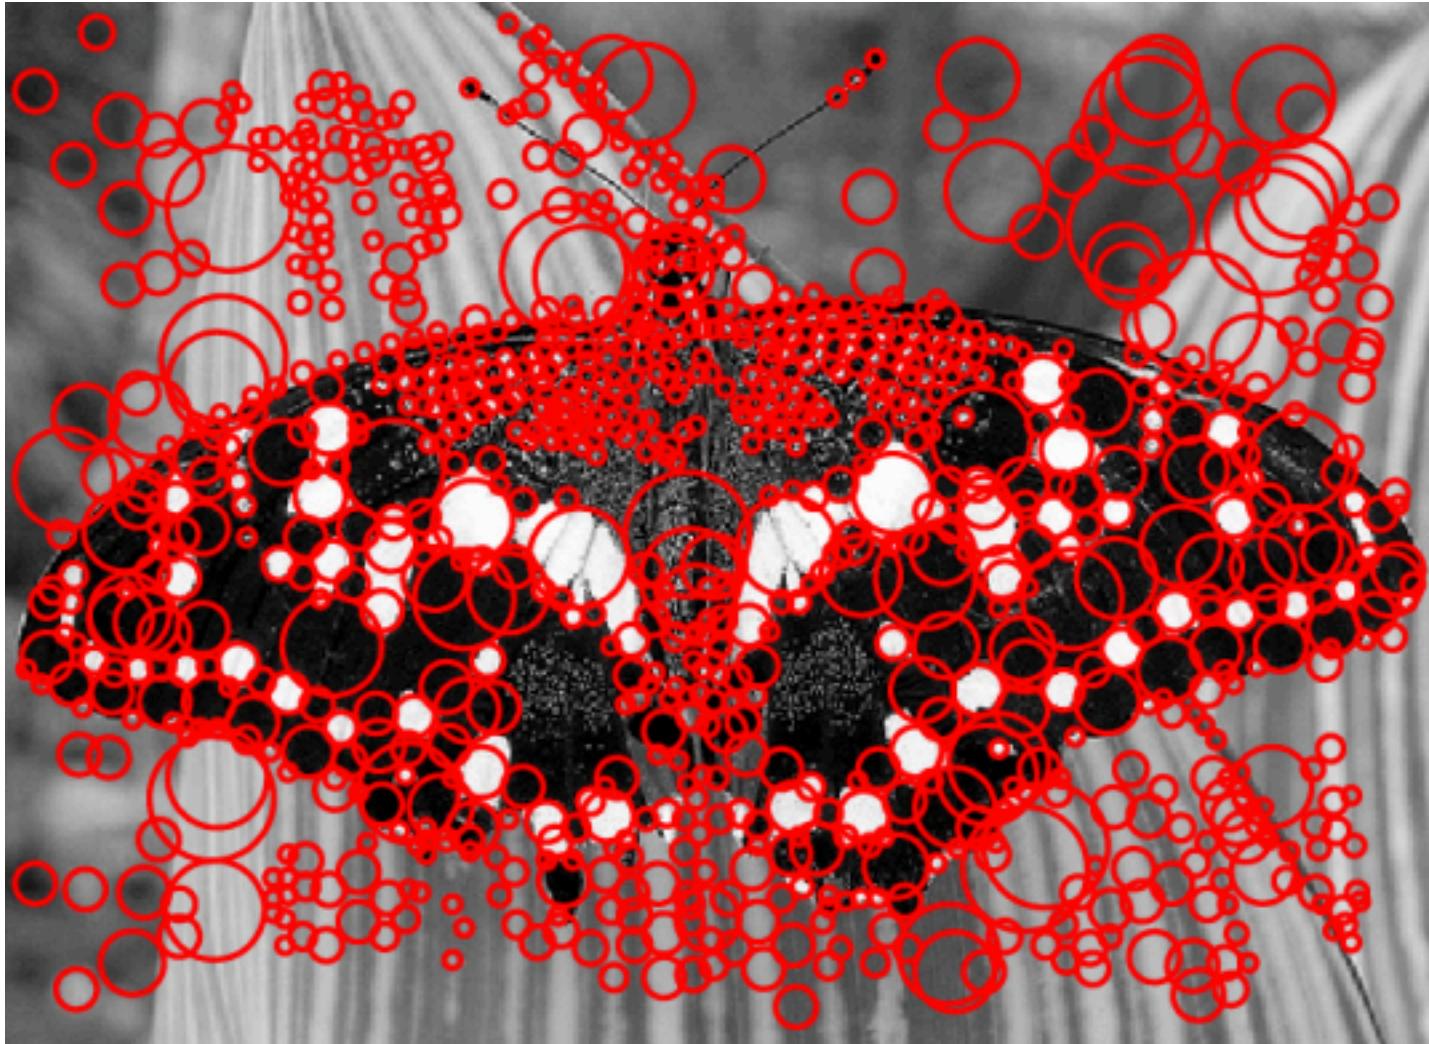
Keypoint Localization

- Detect maxima in scale space
- Then reject points with low image contrast (threshold)
- Eliminate edge responses



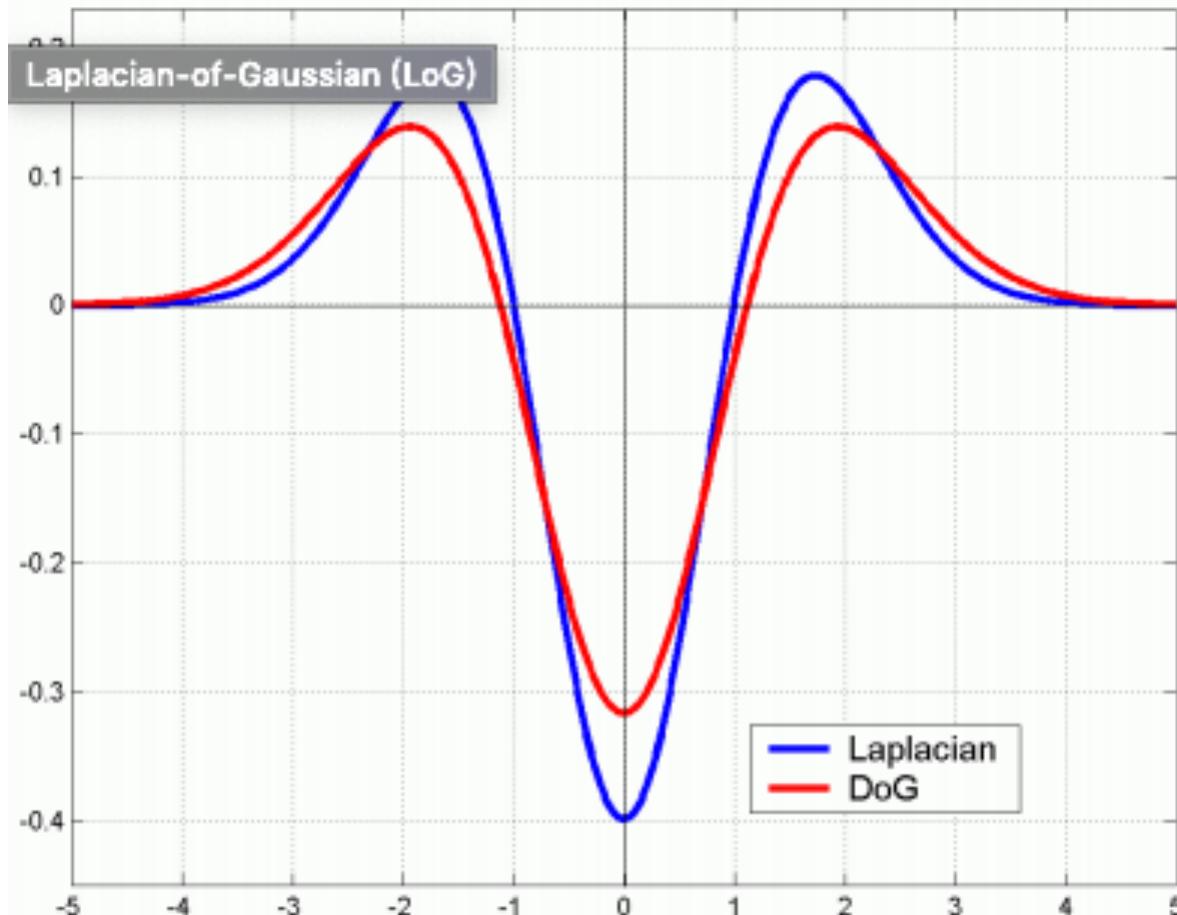
Candidate keypoints: list of (x, y, σ)

LoG Detector



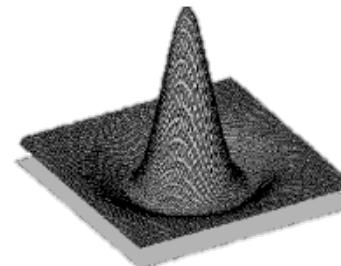
Technical Detail

- We can efficiently approximate the scale-normalized Laplacian with a “difference of Gaussian” (DoG)
 - $L = \sigma^2(G_{xx}(x, y; \sigma) + G_{yy}(x, y; \sigma))$
 - $DoG = G(x, y; k\sigma) - G(x, y; \sigma)$
- Since it is an approximation, DoG is generally slightly inferior to LoG



Difference-of-Gaussians (DoG)

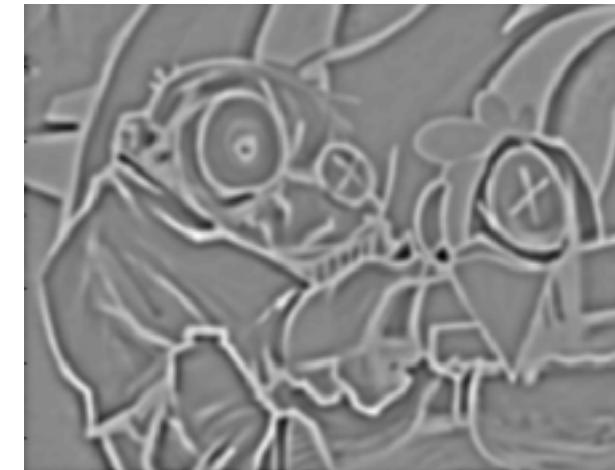
- Difference-of-Gaussians as approximation of the scale-normalized LoG
- Advantage
 - No need to compute 2nd derivatives



-

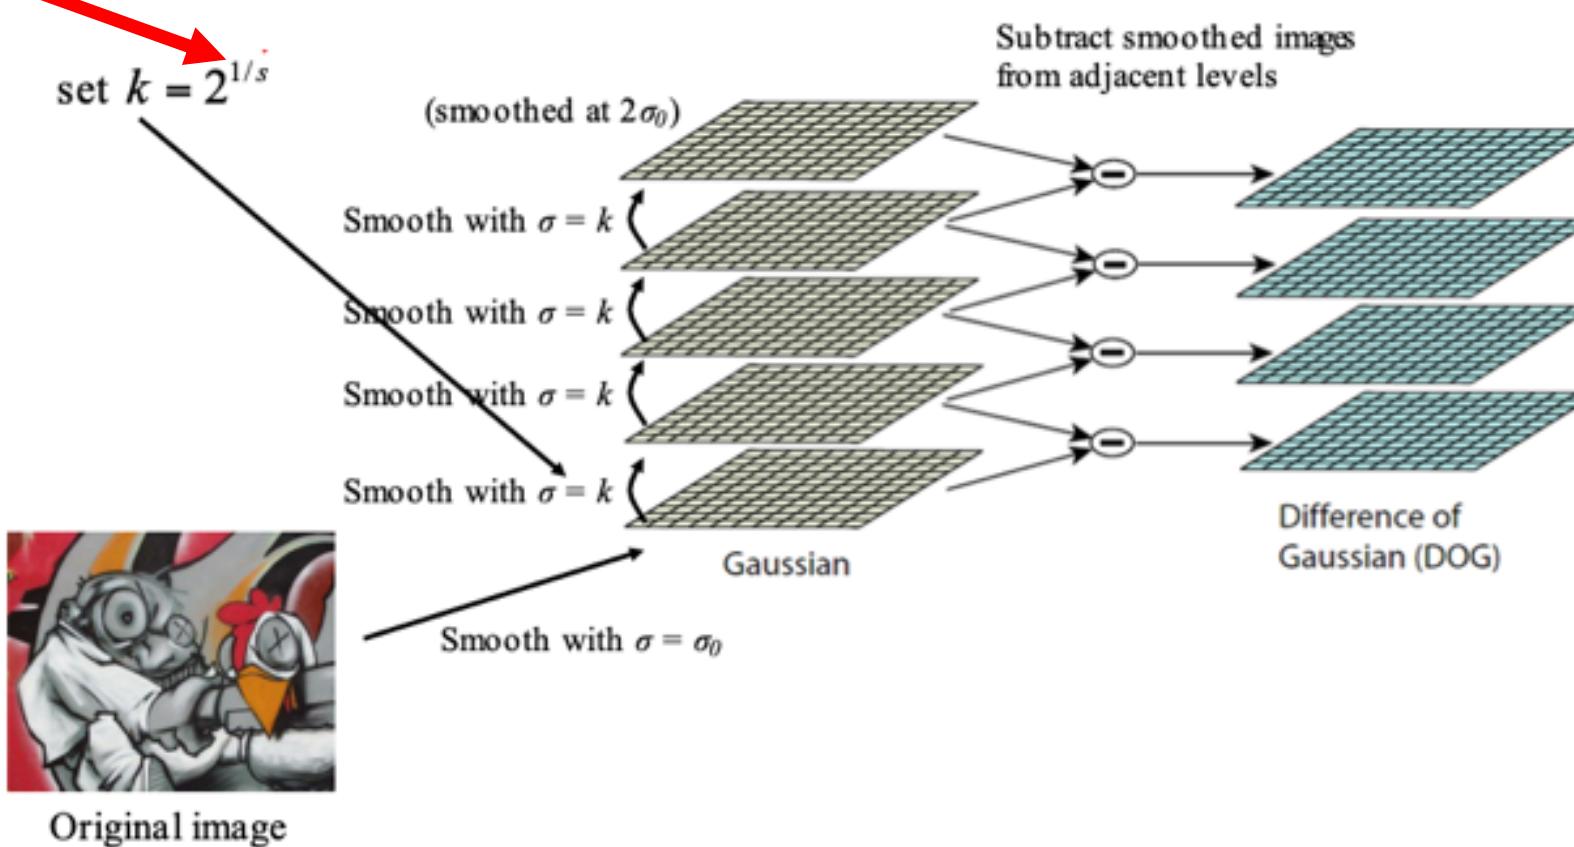


=

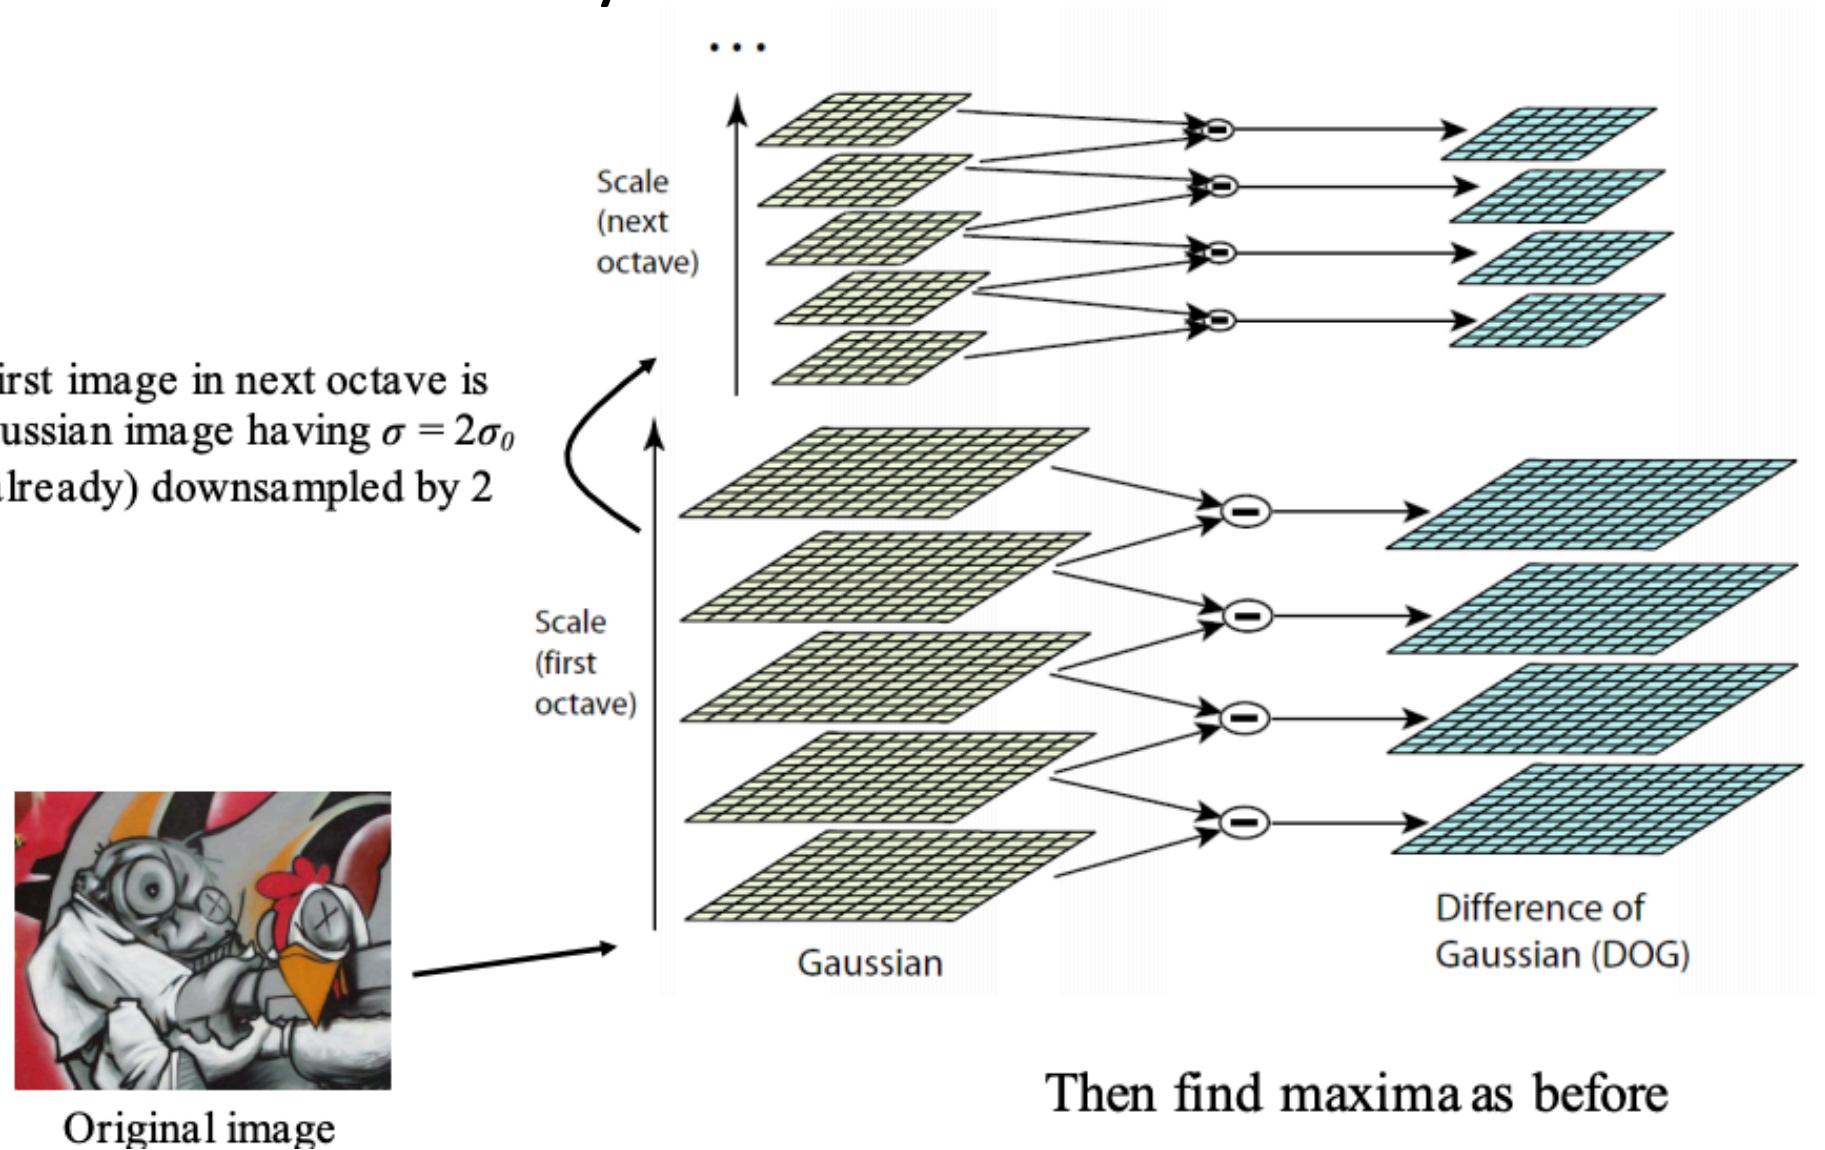


DoG “Octave”

- Octave=“set of images for which smoothing σ is doubled”
- s =desired # of images in DoG “octave”



DoG Octave Pyramid



Example of Keypoint Detection



DoG extrema



Points remaining after
contrast threshold



Points remaining after
removing edge responses

Key Takeaways

- Use a magic function to determine the scale of a feature
 - Laplacian-of-Gaussian (LoG)
- Implementation: use Difference of Gaussian (DoG) to create octave pyramid (approximation)
 - Find the local extreme in scale and space neighbors
- Remove low contrast points and edge points

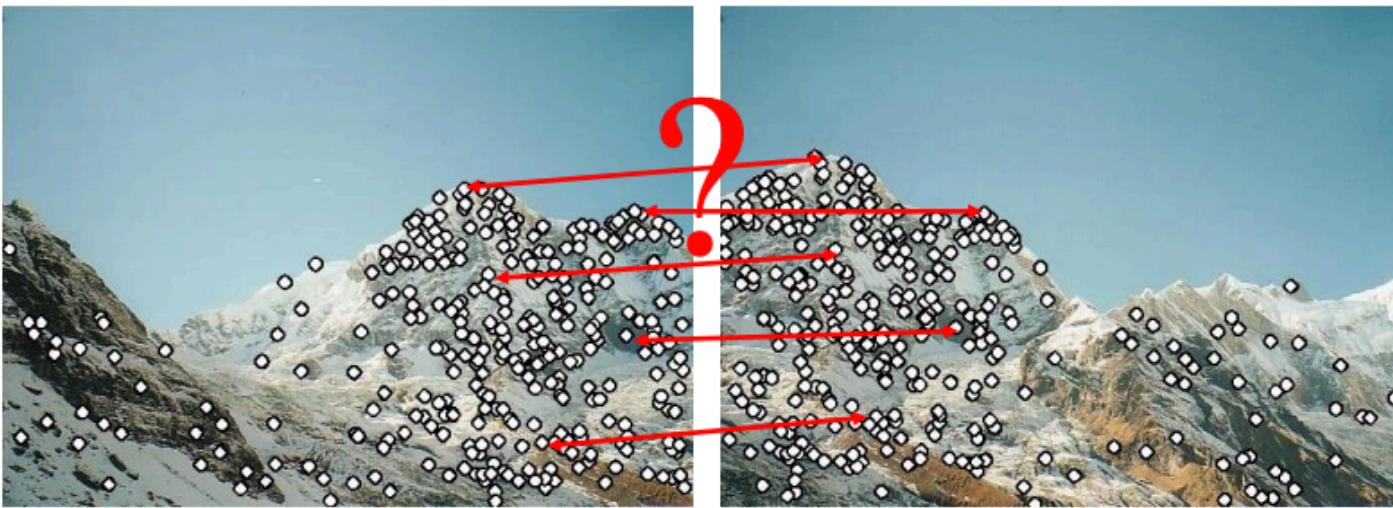
Outline

- Local Invariant Features
 - Motivation
 - Requirements, Invariances
- Keypoint Localization
 - Feature from Accelerated Segment Test (FAST)
 - Harris
 - Shi-Tomasi
- Scale Invariant Region Selection
 - Automatic scale selection
 - Laplacian-of-Gaussian detector
 - Difference-of-Gaussians detector
- Local Descriptors
 - Orientation normalization
 - SIFT

Local Descriptors

- We know how to detect points and appropriate scales
- Next question:

How to describe them for matching?



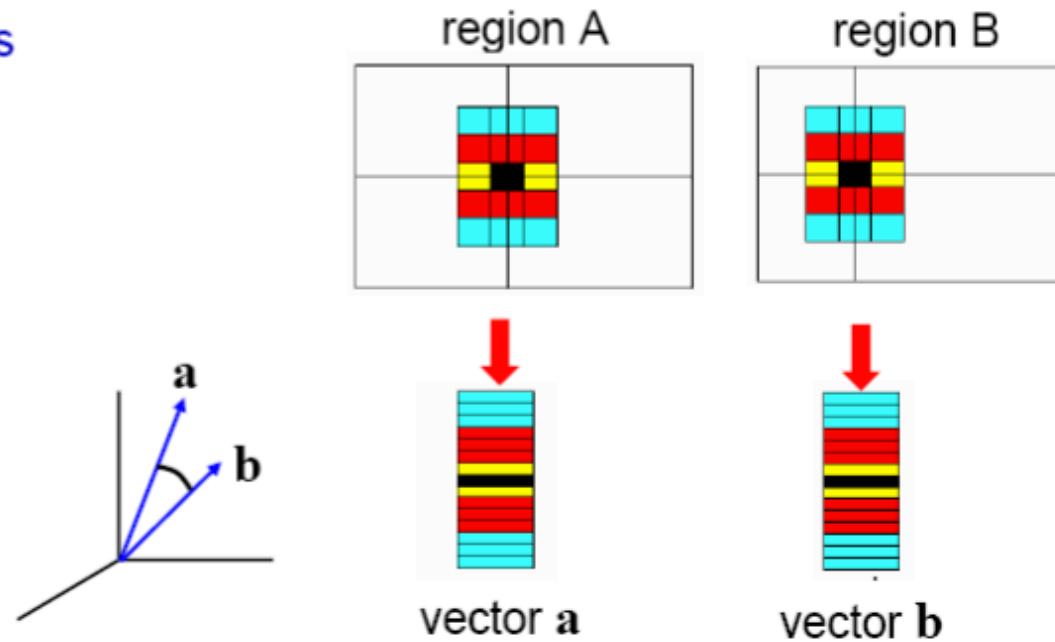
Point descriptor should be:
1. Invariant
2. Distinctive

Local Descriptors

- Simplest descriptor: rasterized list of intensities within a patch
- To what is this going to be invariant (for normalized patch sizes)?
 - Translation and scale only (not rotation)

Write regions as vectors

$$A \rightarrow \mathbf{a}, \quad B \rightarrow \mathbf{b}$$

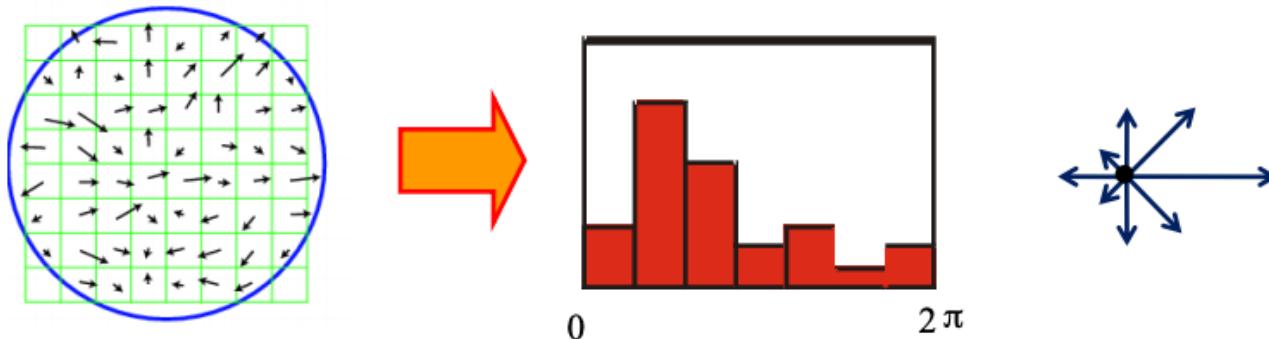


Feature Descriptors

- Disadvantage of patches as descriptors:
 - Small deformations can greatly affect matching score

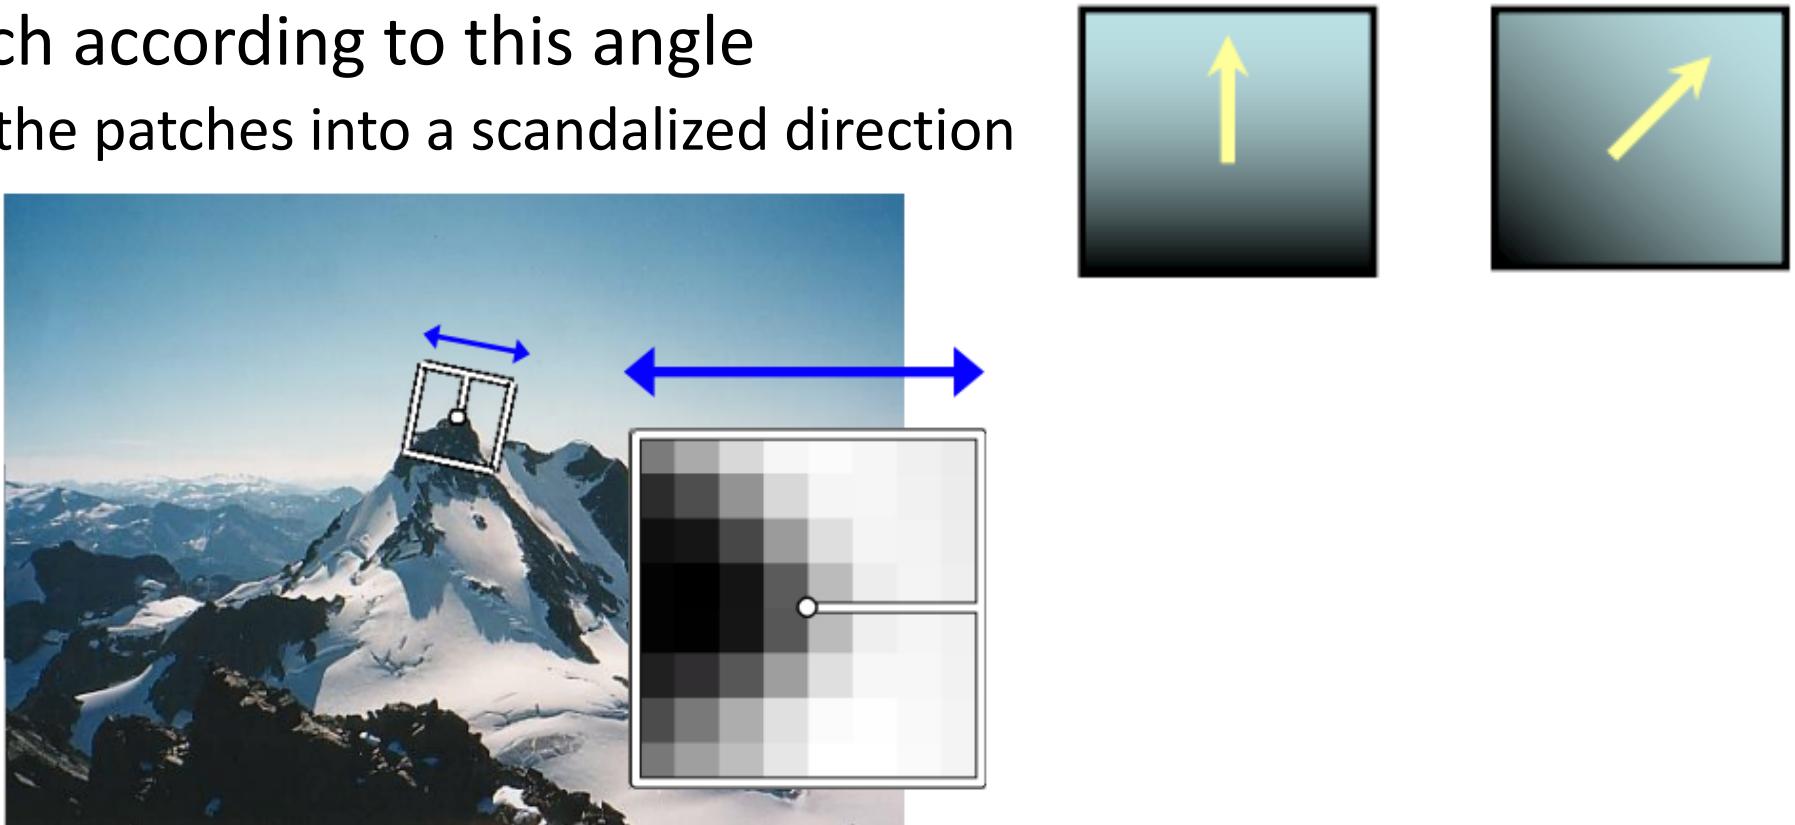


- Solution: histogram of gradient direction



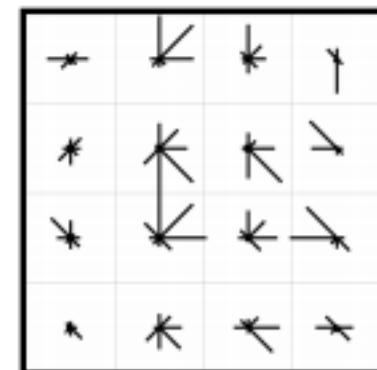
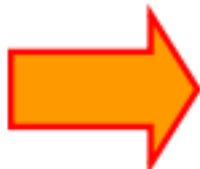
Rotation Invariant Descriptors

- Find local gradient directions in patch
- Find dominant direction of overall gradient for the patch
- Rotate patch according to this angle
 - This puts the patches into a standardized direction



Feature Descriptors: SIFT

- **Scale Invariant Feature Transform**
- Descriptor computation:
 - Divide patch (for keypoint at chosen scale) after standardized direction alignment into 4x4 sub-patches: 16 cells
 - Compute Gaussian weighted histogram of gradient directions (8 reference angles) for all pixels inside each sub-patch
 - Resulting descriptor: $4 \times 4 \times 8 = 128$ dimensions
 - Rasterize and normalize to unit vector (enhances invariance to affine changes in illumination)
 - Clip high values and renormalize vector to reduce effects of non-linear illumination



Working with SIFT Descriptors

- One image with n feature point yields:
 - n 128-dimensional descriptor: each one is a collection of histograms of the gradients orientations within a patch
 - [$n \times 128$ matrix]
 - n scale parameters specifying the size of each patch
 - [$n \times 1$ matrix]
 - n orientation parameters specifying the angle of the patch
 - [$n \times 1$ matrix]
 - n 2-D points giving positions of the patches
 - [$n \times 2$ matrix]

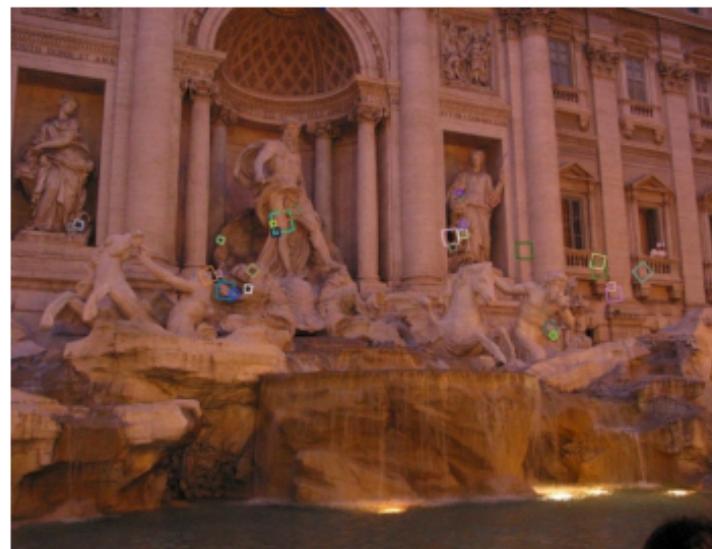


Matching SIFT Descriptors

- Given two images
- Match each keypoint independently
 - Employ 128-D descriptors
 - Best candidate match is closest Euclidean neighbor
 - Discard matches using ratio of closest and second closest neighbors
 - Simple measure of how unique is the match
- Can employ more costly approaches to match complete sets of points (e.g., Hungarian algorithm)

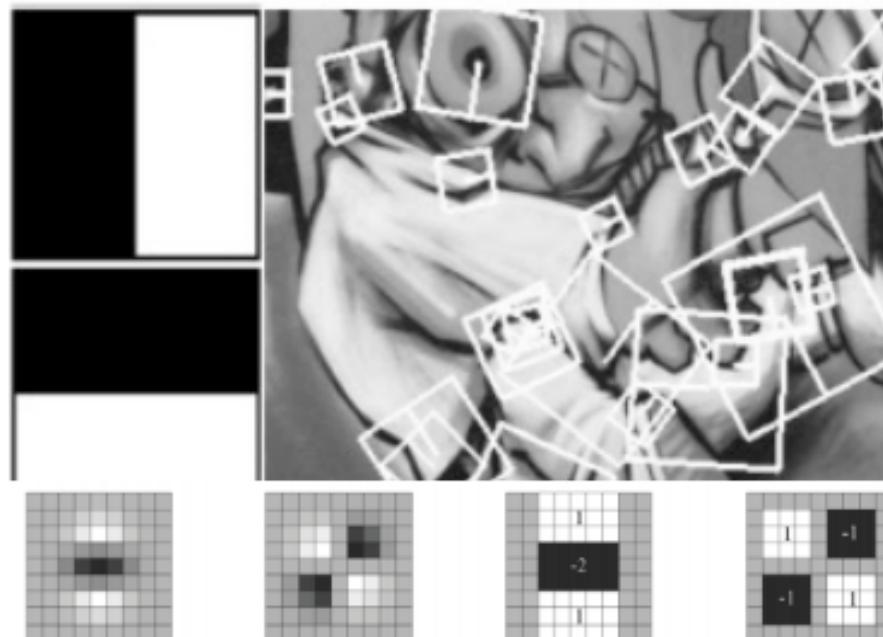
Advantages of SIFT

- Fairly robust matching technique
 - Can handle significant change in viewpoint
 - Can handle significant changes in illumination
 - Sometimes even day vs night
 - Fast and efficient – can run in real time
 - Lots of code available



Variations and Implementations

- SURF: Fast approximation of SIFT (multiple times faster)
 - Efficient computation by 2-D box filters and integral images
 - Essentially equivalent quality for object identification
- GPU implementations of SIFT and SURF exist



Key Takeaways

- **SIFT** is a very popular and powerful local descriptor from 15-20 years ago
 - Translation, rotation and scale invariant
 - Handle significant illumination and viewpoint change well
- Translation invariant:
- Rotation invariant: define/find the dominant direction of overall gradient for the patch
- Scale: use “Octave” to define the scale of a patch
- A lot of open source code and various SIFT implementations